
# Elastic Stack Reference Architecture

**Elastic**
**www.elastic.co**

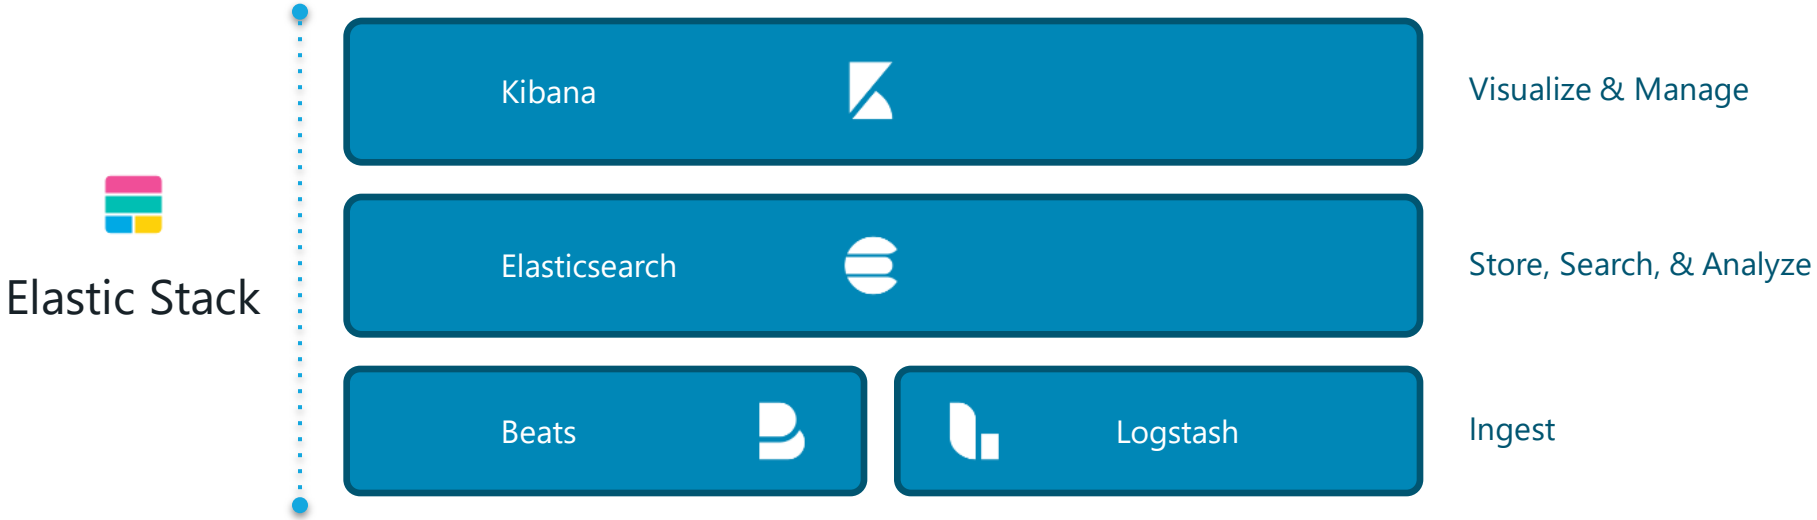
Elastic Stack
Kibana
Visualize & Manage
Elasticsearch
Store, Search, & Analyze
Beats
Logstash
Ingest



elastic

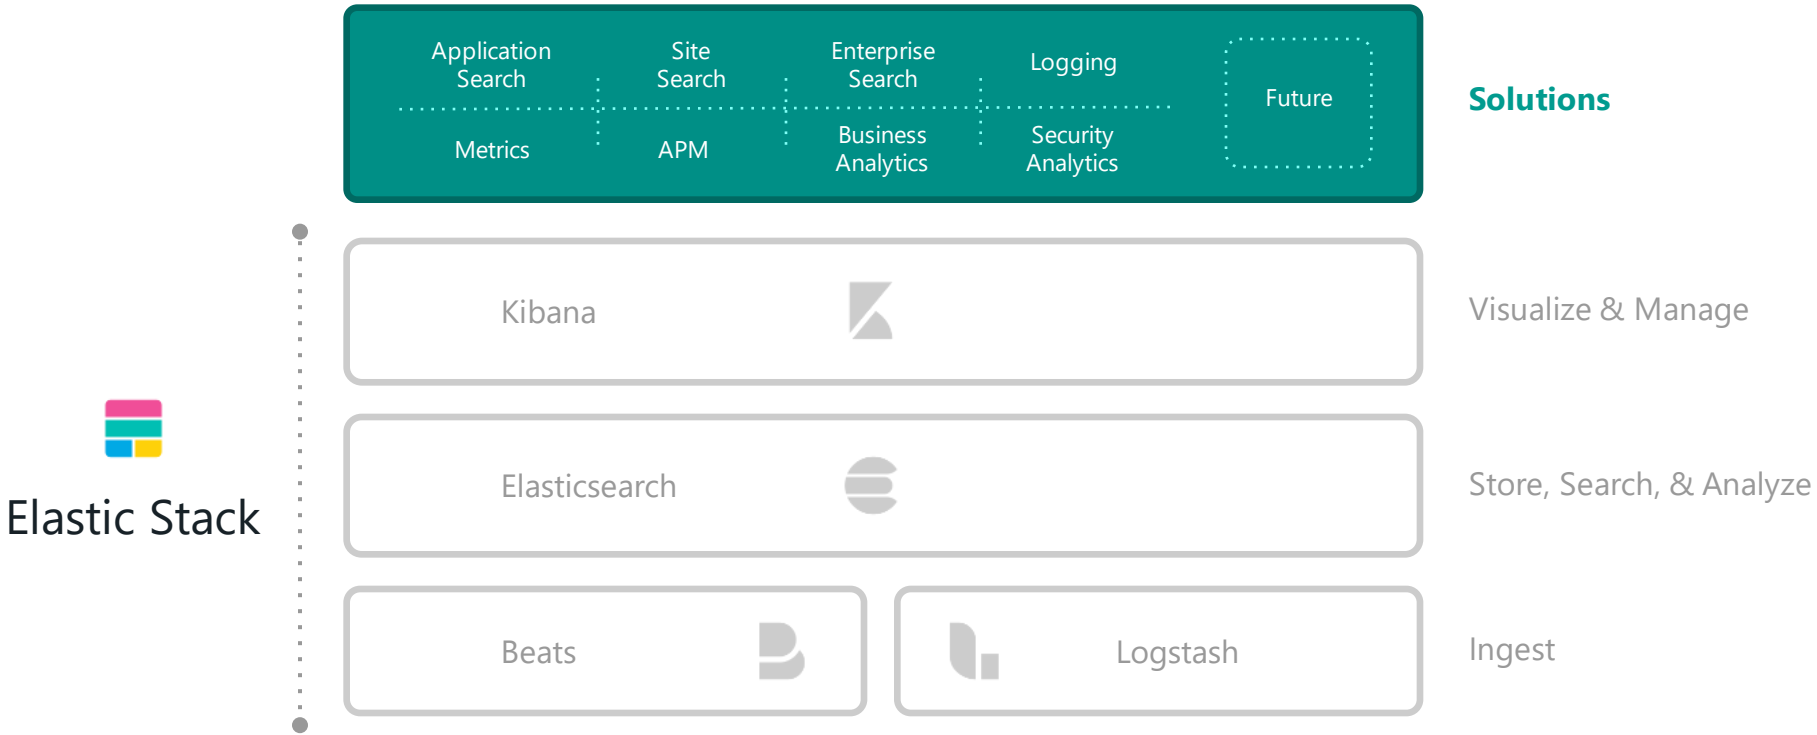
Application Search
Site Search
Enterprise Search
Logging
Future
Metrics
APM
Business Analytics
Security Analytics
Solutions
Kibana
Visualize & Manage
Elasticsearch
Store, Search, & Analyze
Beats
Logstash
Ingest
Elastic Stack



elastic

# Solutions

APM

Site Search

App Search

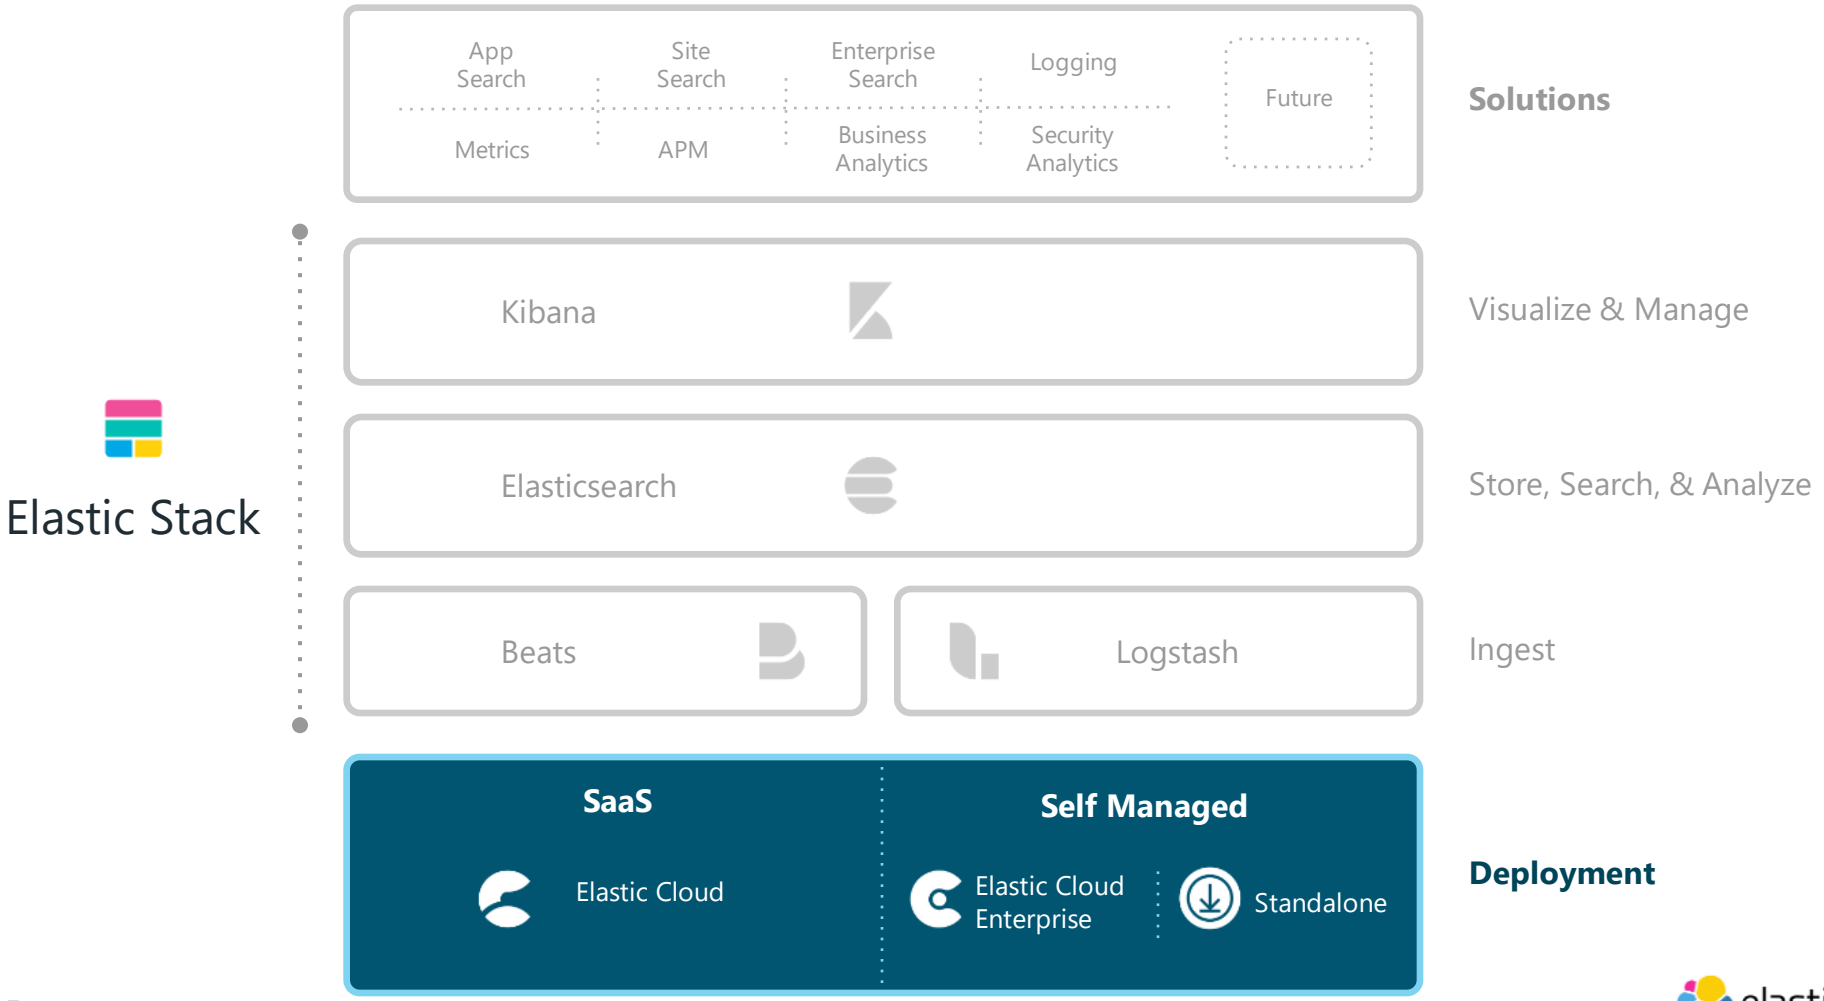
App Search
Site Search
Enterprise Search
Logging
Metrics
APM
Business Analytics
Security Analytics
Future
Solutions
Elastic Stack
Kibana
Visualize & Manage
Elasticsearch
Store, Search, & Analyze
Beats
Logstash
Ingest
SaaS
Elastic Cloud
Self Managed
Elastic Cloud Enterprise
Standalone
Deployment

# SaaS

Elastic Cloud

Elasticsearch Service

Site Search

App Search

# Self Managed

Elastic Cloud Enterprise

Standalone



elastic

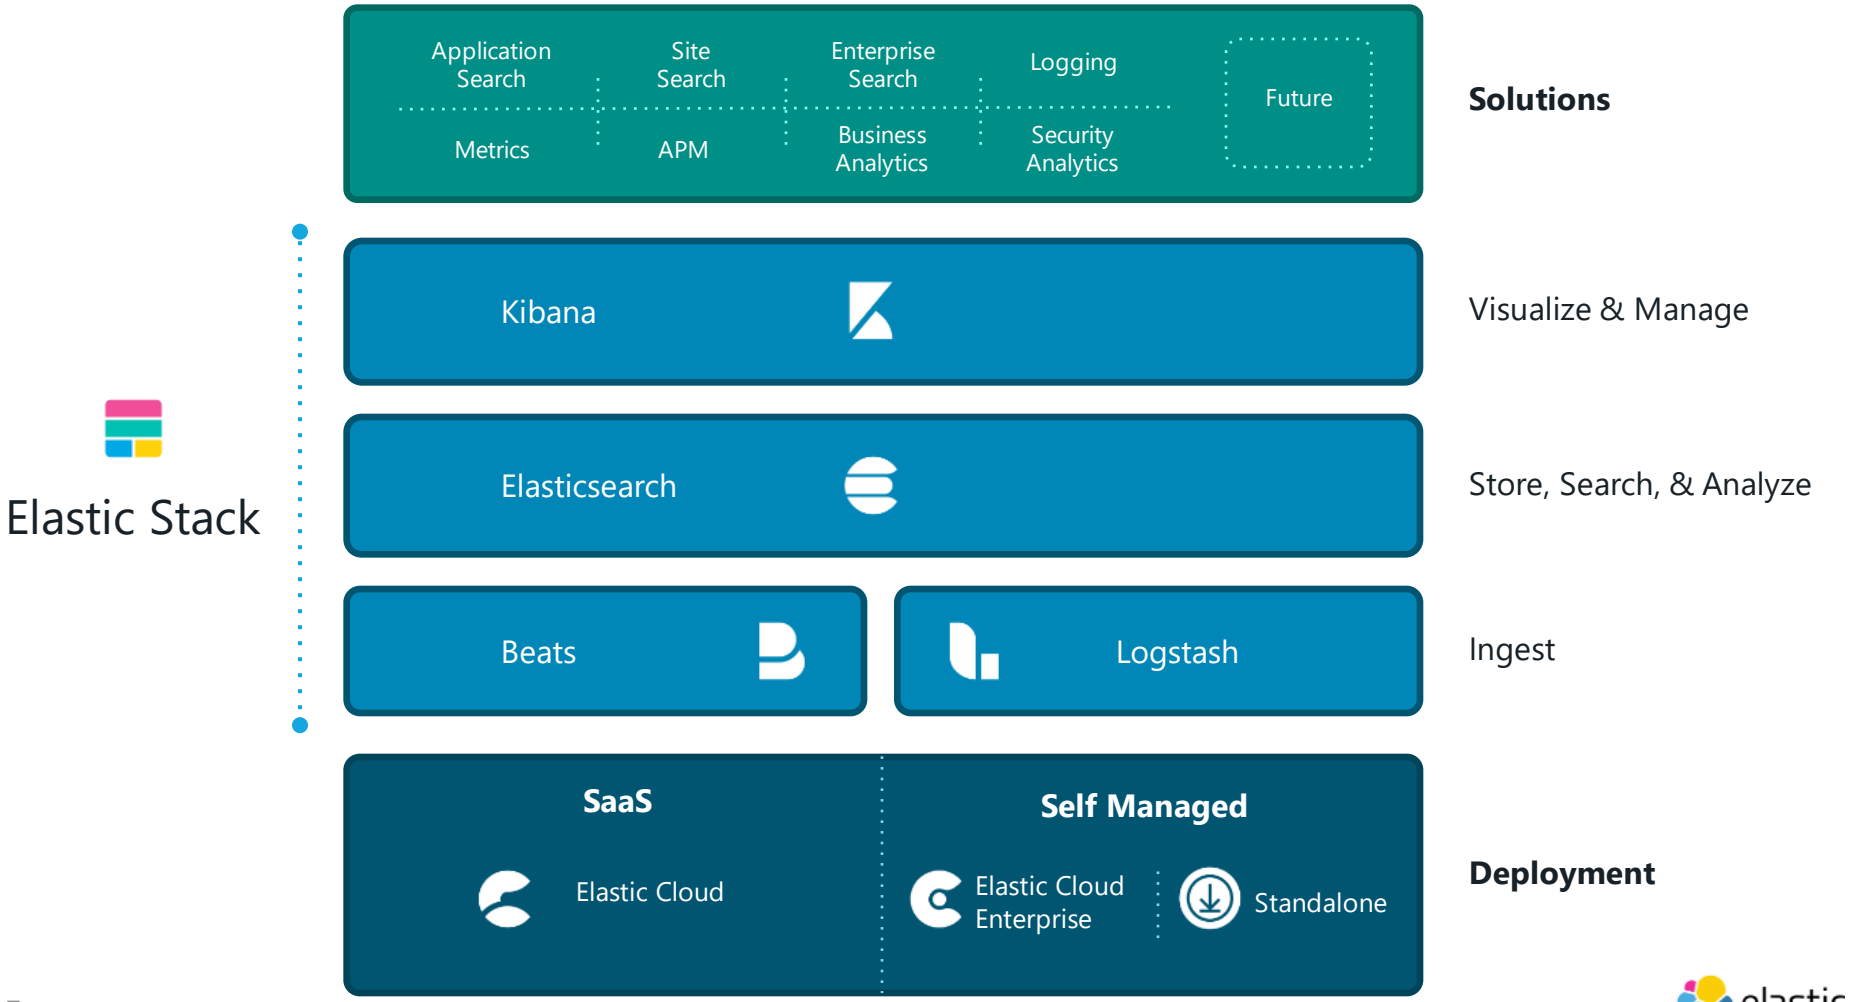
Application Search
Site Search
Enterprise Search
Logging
Future
Metrics
APM
Business Analytics
Security Analytics
Solutions
Kibana
Visualize & Manage
Elasticsearch
Store, Search, & Analyze
Beats
Logstash
Ingest
Elastic Stack
SaaS
Elastic Cloud
Self Managed
Elastic Cloud Enterprise
Standalone
Deployment

# Logical Processing Pipeline

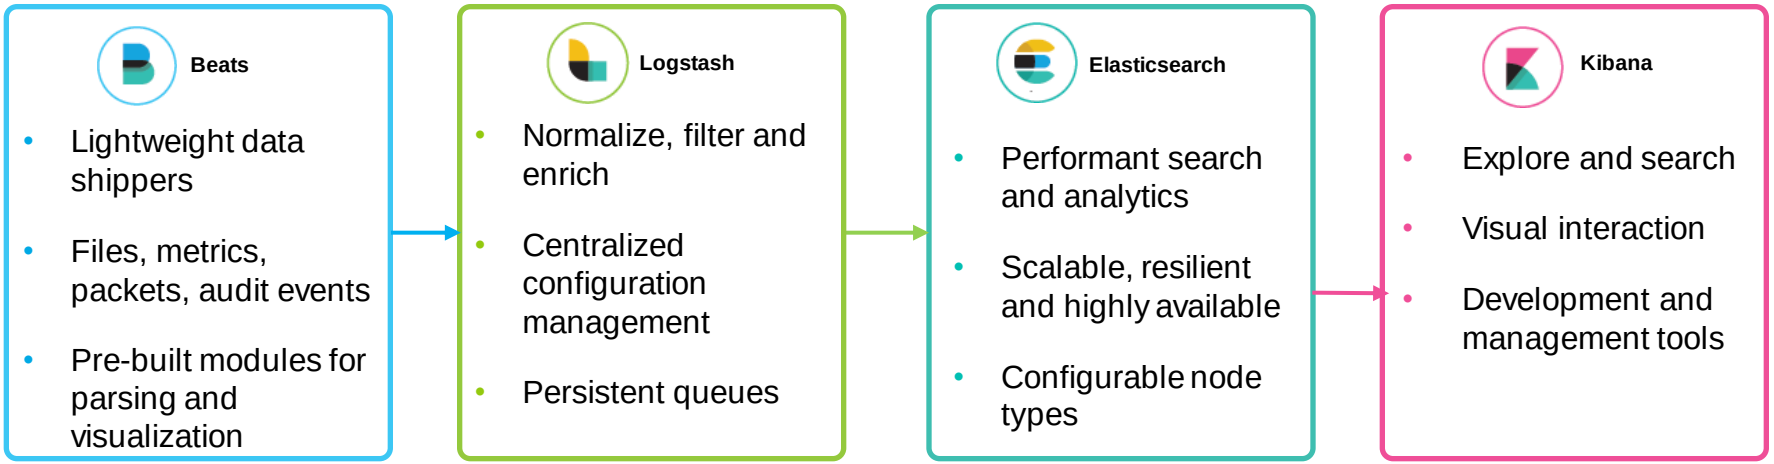






elastic

# Logging Architecture

# Basic Processing Pipeline

*Beats, Elasticsearch w/ Ingest Node Pipelines and Kibana*

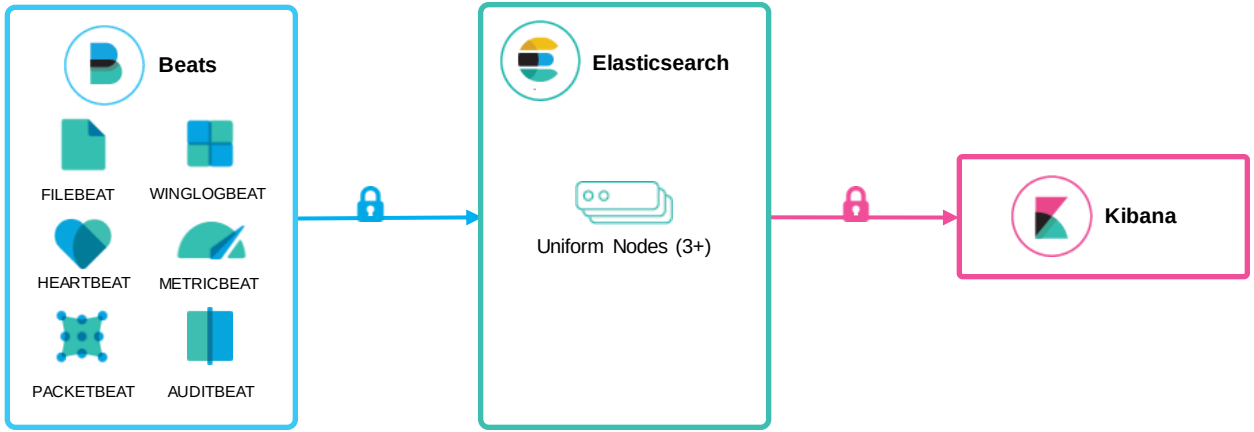

# Advanced Processing and Resiliency

*Adding Logstash processing, differentiated Elasticsearch node types*

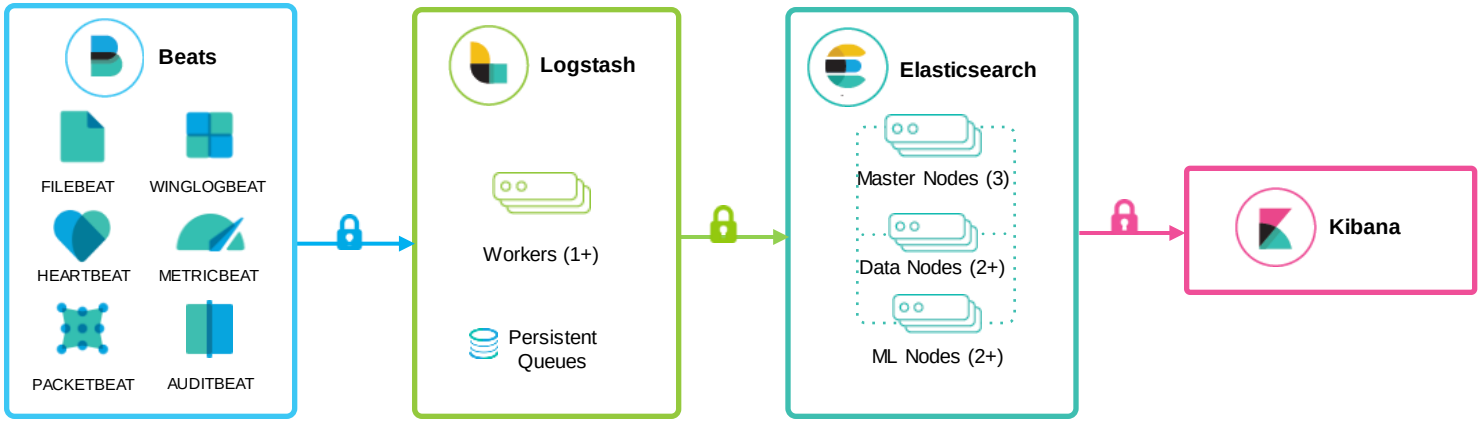




elastic

# Elastic Cloud Based Logging

*Capturing events securely to Elastic Cloud with easy configuration*

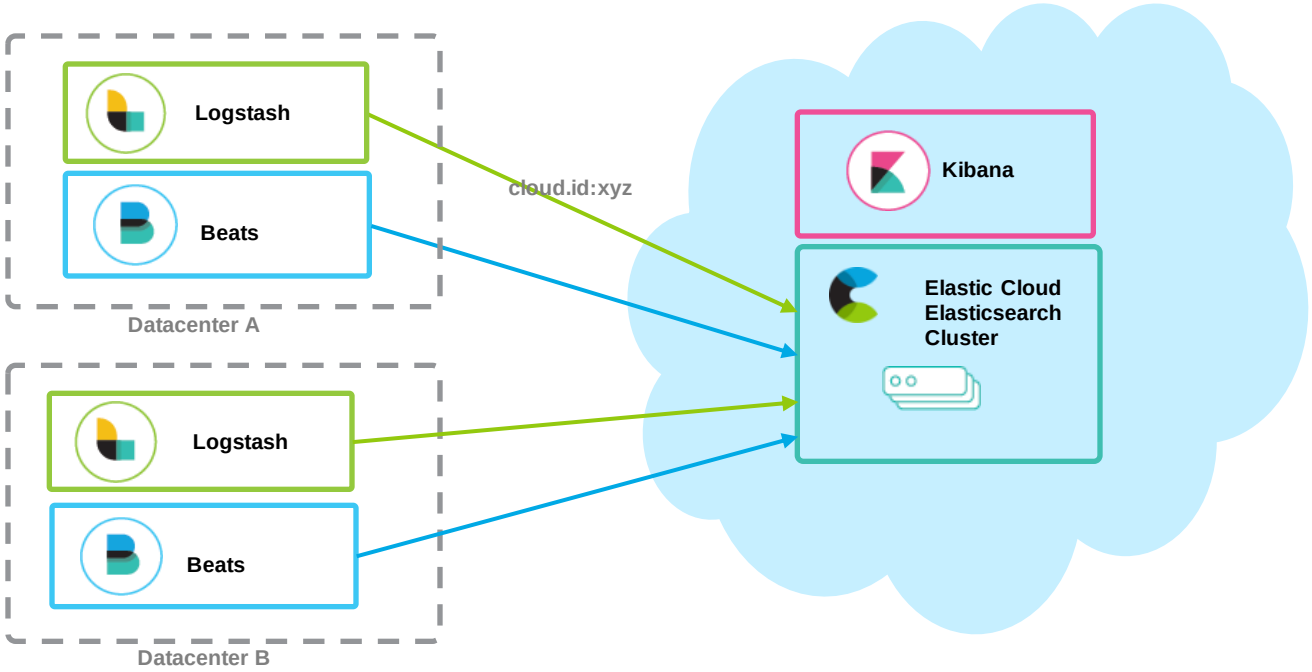

# Deployment in the Enterprise

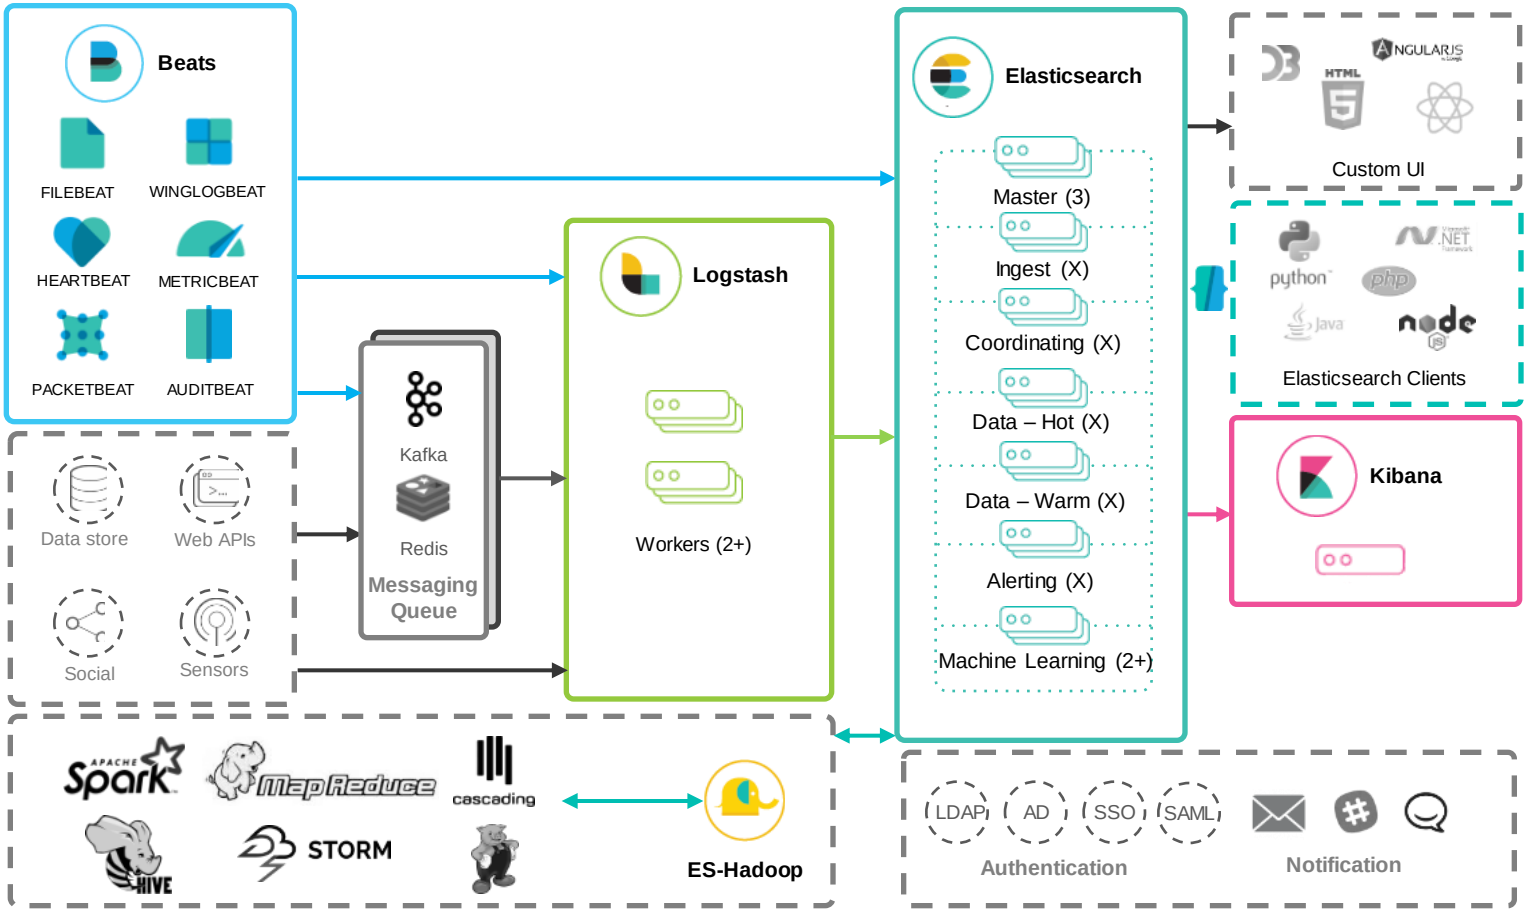






elastic

# Options for Migrating from Self-Managed to Elastic Cloud

*Multiple options for easily migrating to Elastic Cloud from self-managed*

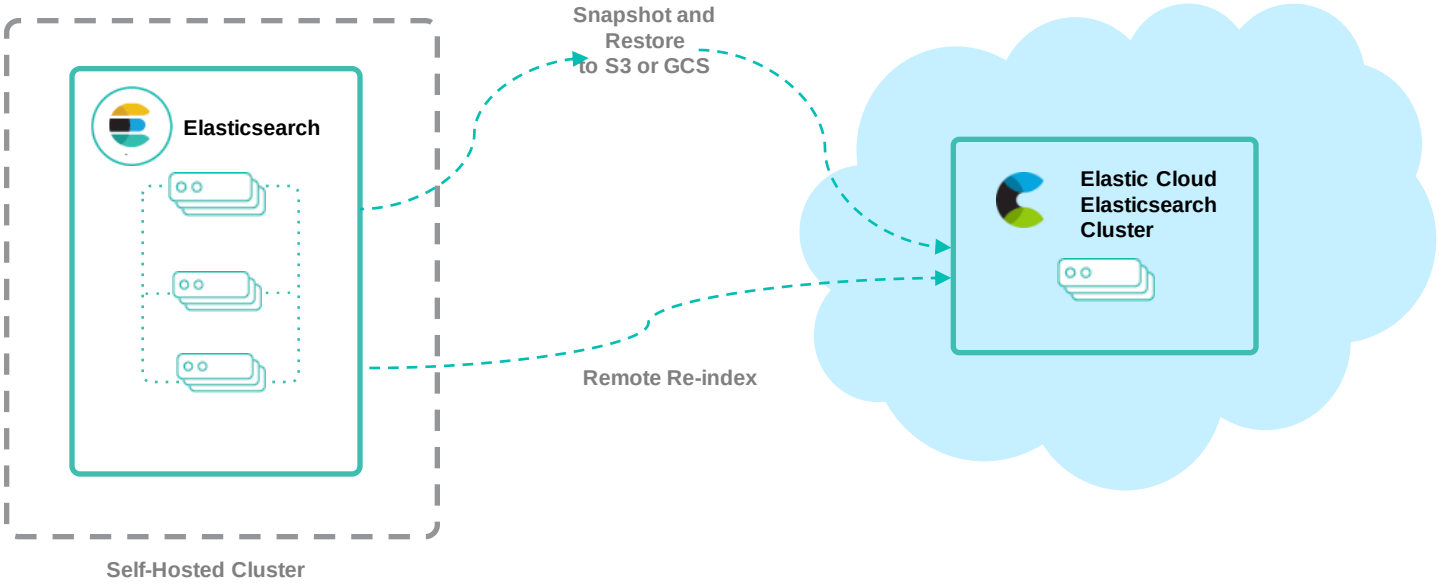

# Inside a Large Elasticsearch Logging Cluster

*Reduce infrastructure costs, isolate workloads, and manage data lifecycle*

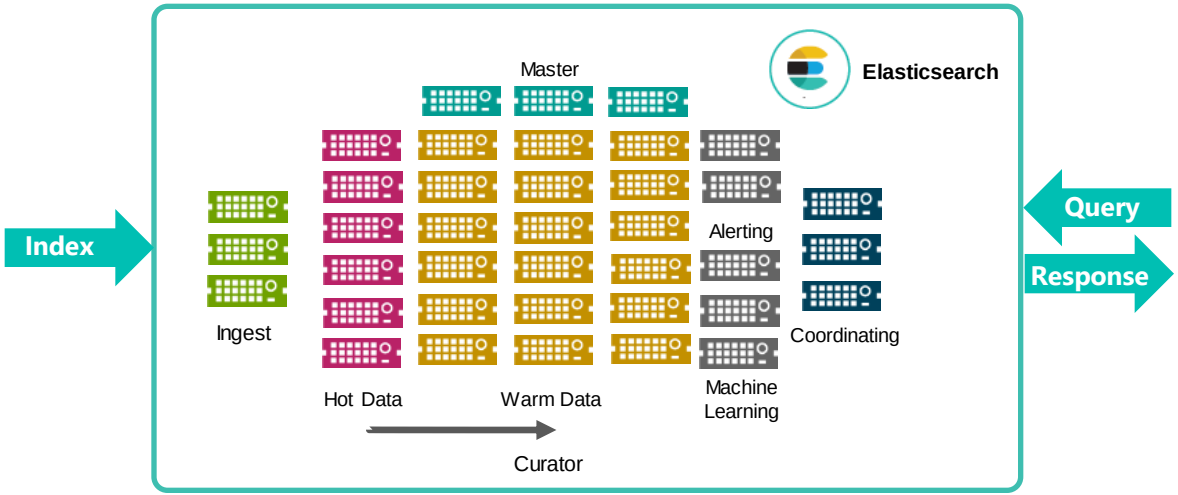

# Variants of Logging Architectures

# Inside Logstash

*Example Inputs, Filters and Outputs*

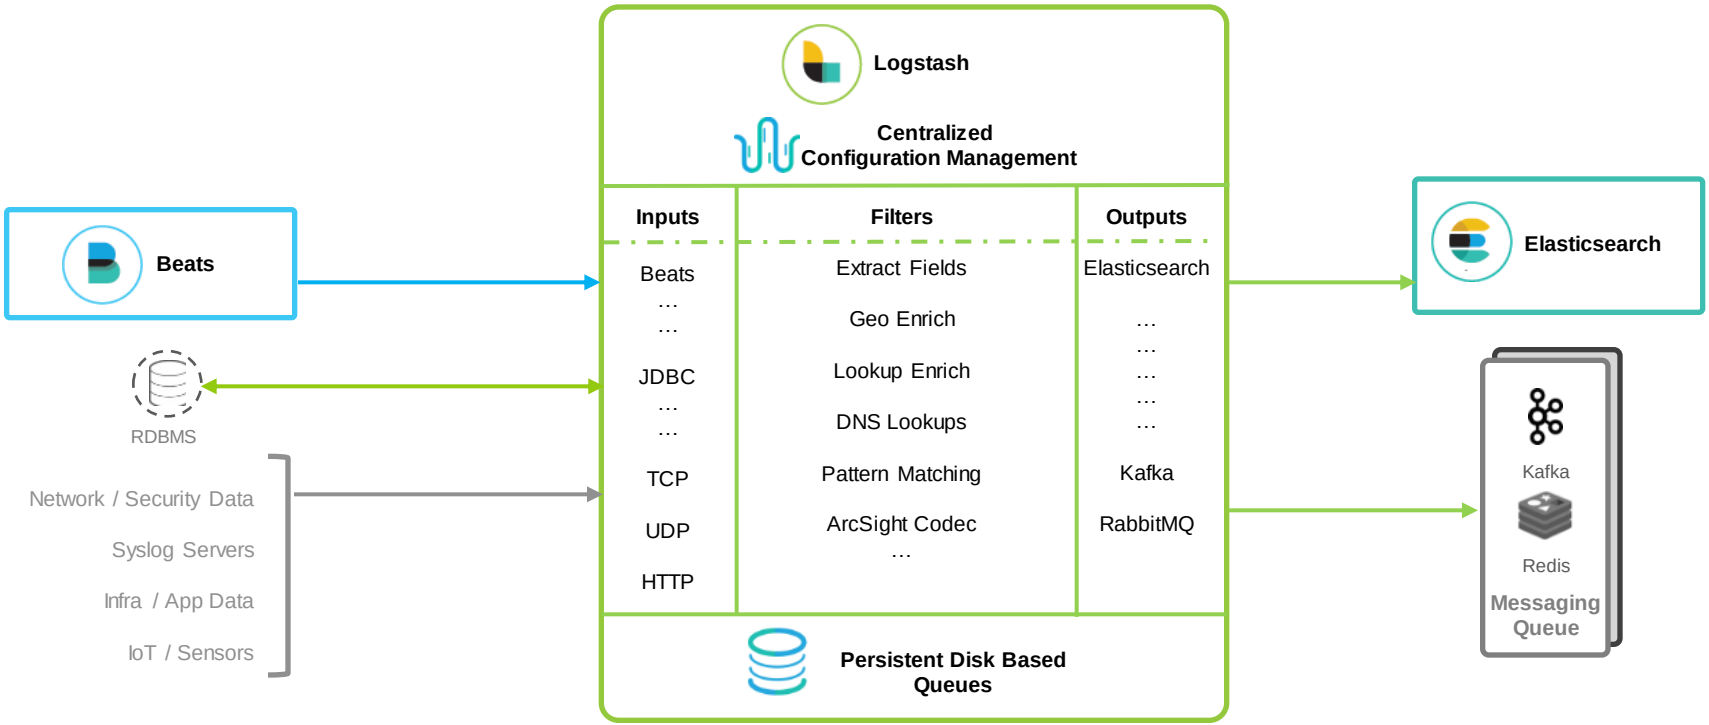




elastic

# At-least-once Delivery with Logstash

*Using inputs and outputs with acknowledgement capability*

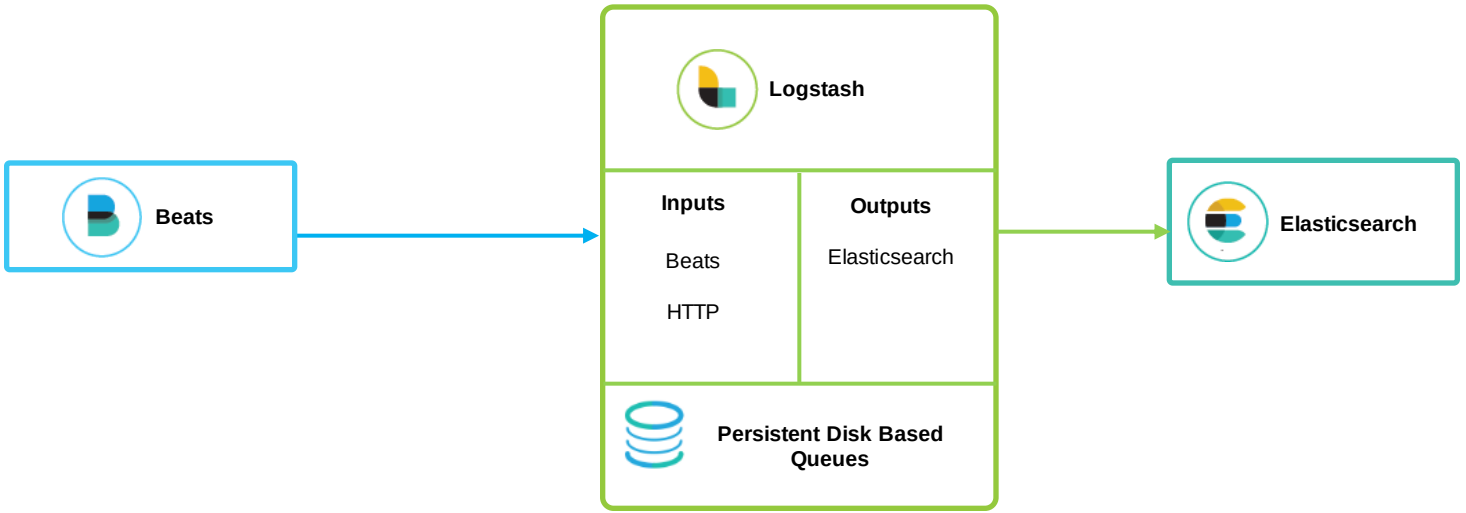




elastic

# Highly Available, Persistent Logstash without an MQ

*Round-robin with persistent Logstash queues responds to backpressure and continues operation with AZ interruption.*

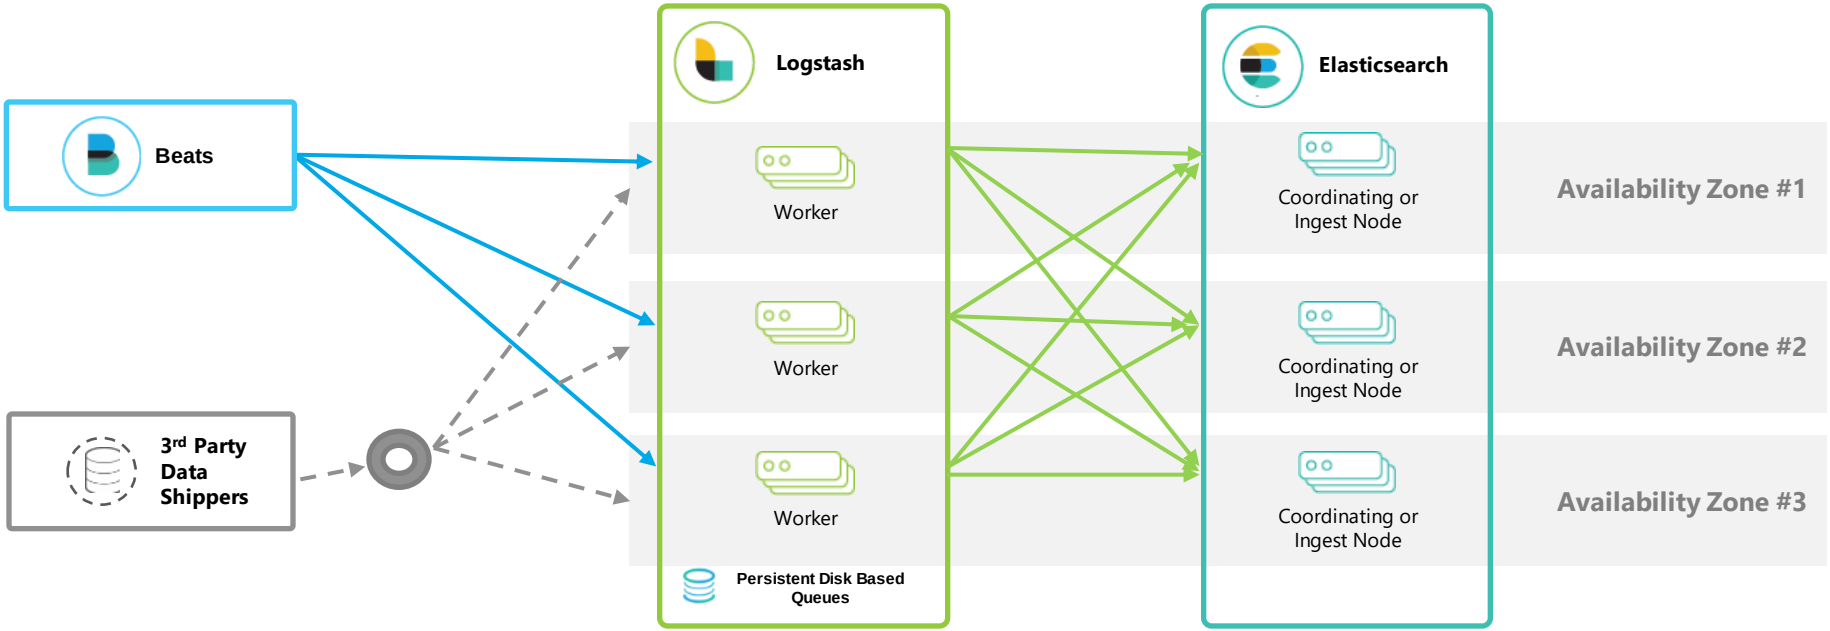






elastic

# Highly Available, Persistent Logstash with an MQ

*Adding a message queue allows for replay and queue replication, thereby avoiding message loss in the case of a unrecoverable Logstash disk failure.*

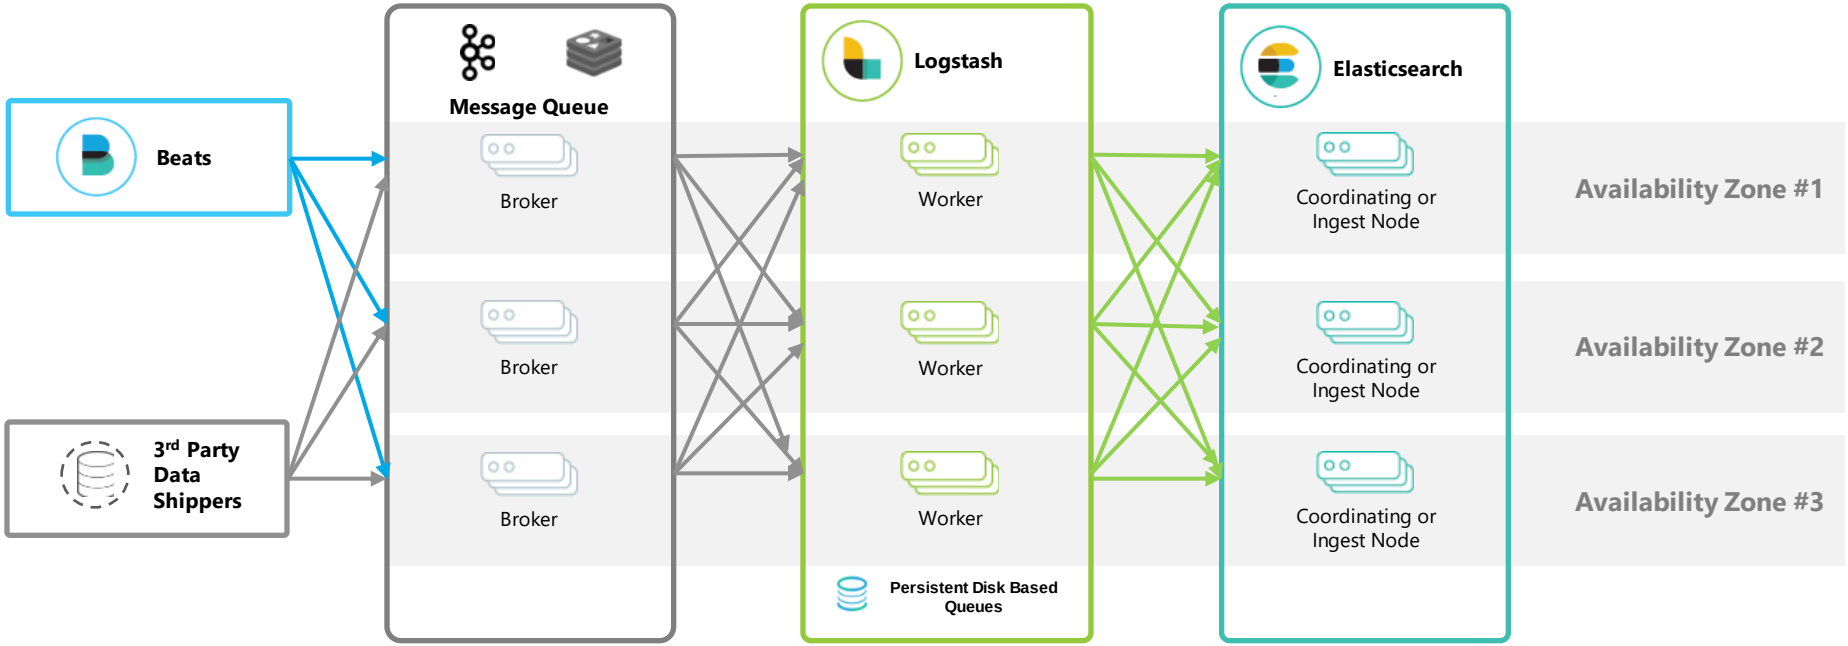

# Control of Logstash Configuration by Group

*Distribute Logstash configuration management while controlling central pipeline*

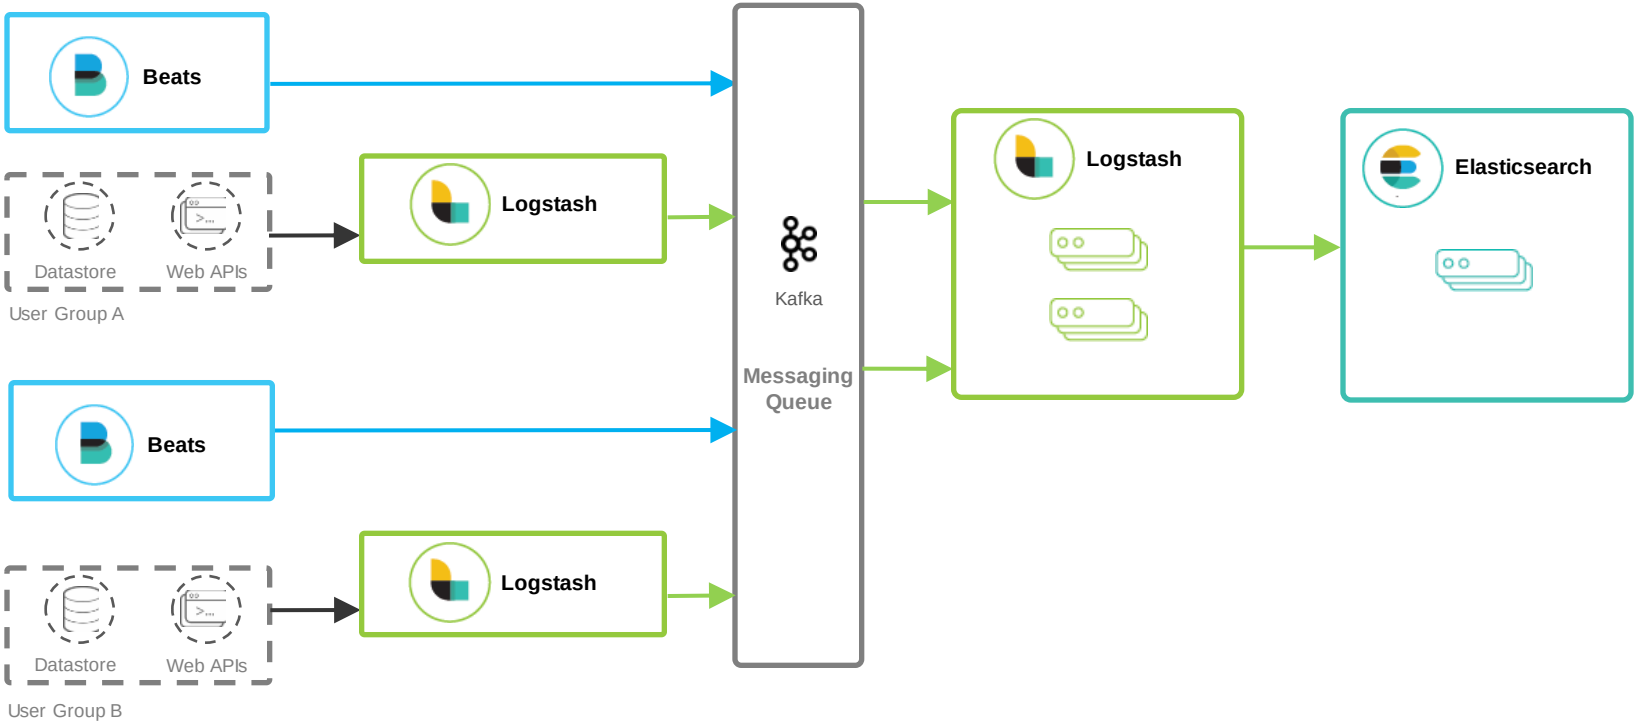




elastic

# Deployment Best Practices

# Centralized Monitoring Cluster

*Maintain isolated monitoring cluster for monitoring workload isolation*

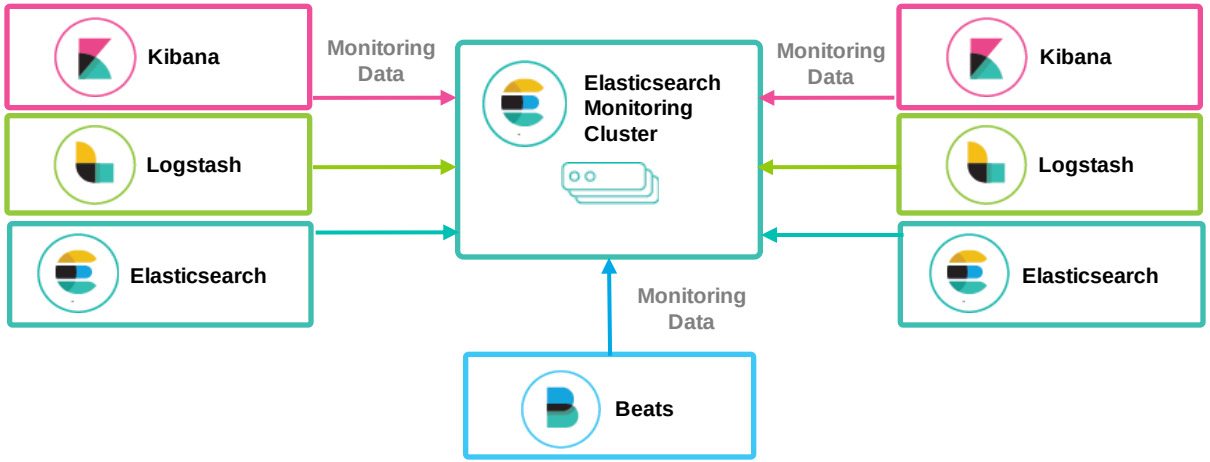

# Cloud Monitoring Cluster

*Opt-in Elastic Cloud cluster for monitoring on-premise stack*

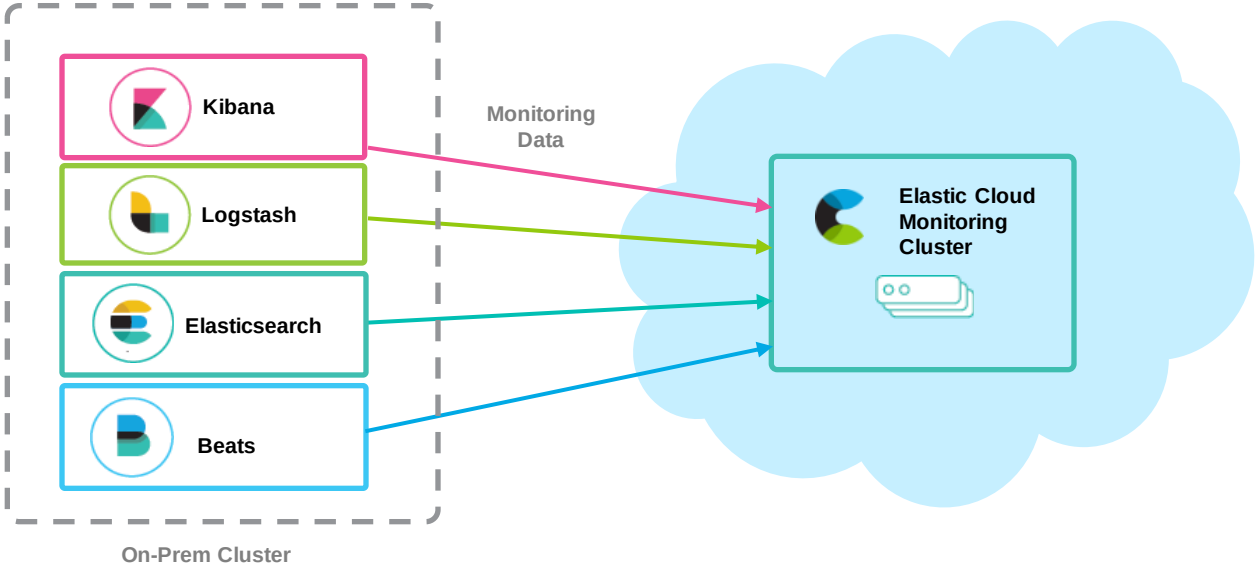

# Isolated Audit Logging Cluster

*Maintain isolated audit logging cluster for increased security and compliance*

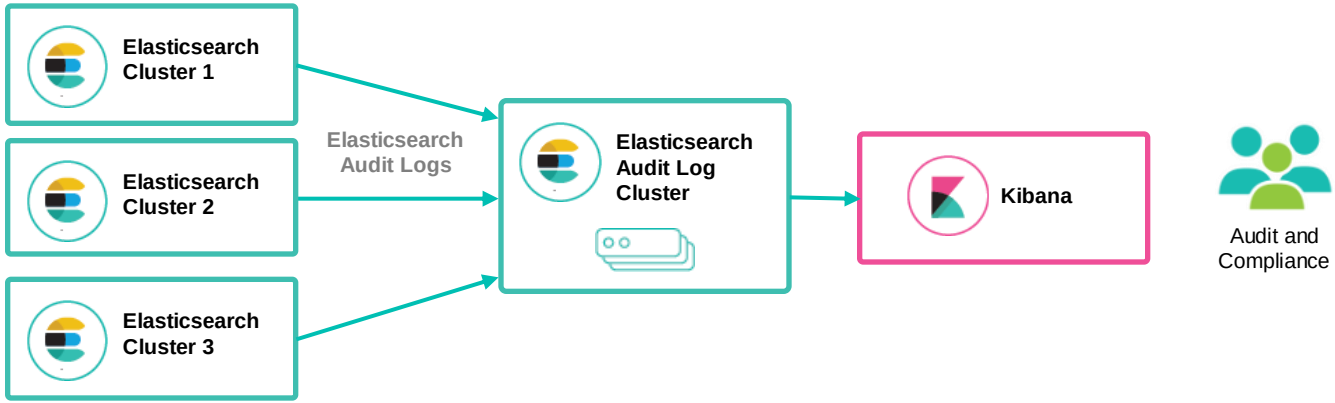

# Use-Case Architectures

# Application Metric Collection with Elastic APM

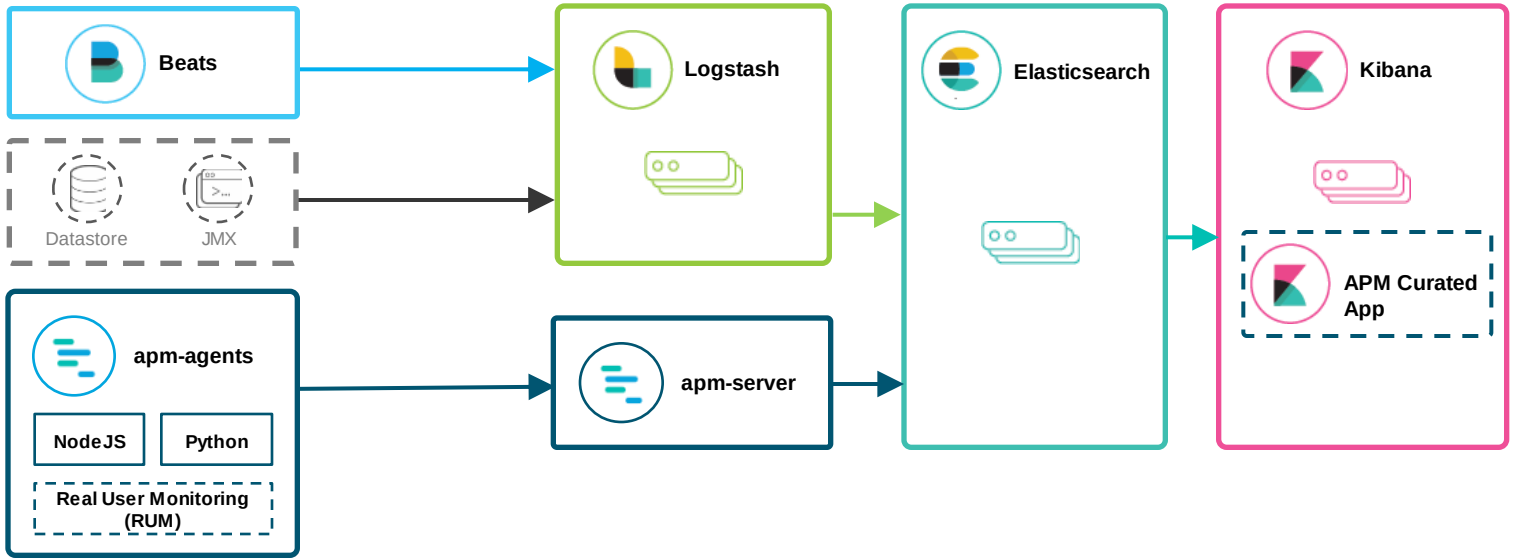




elastic

# Application or Enterprise Search: Compact

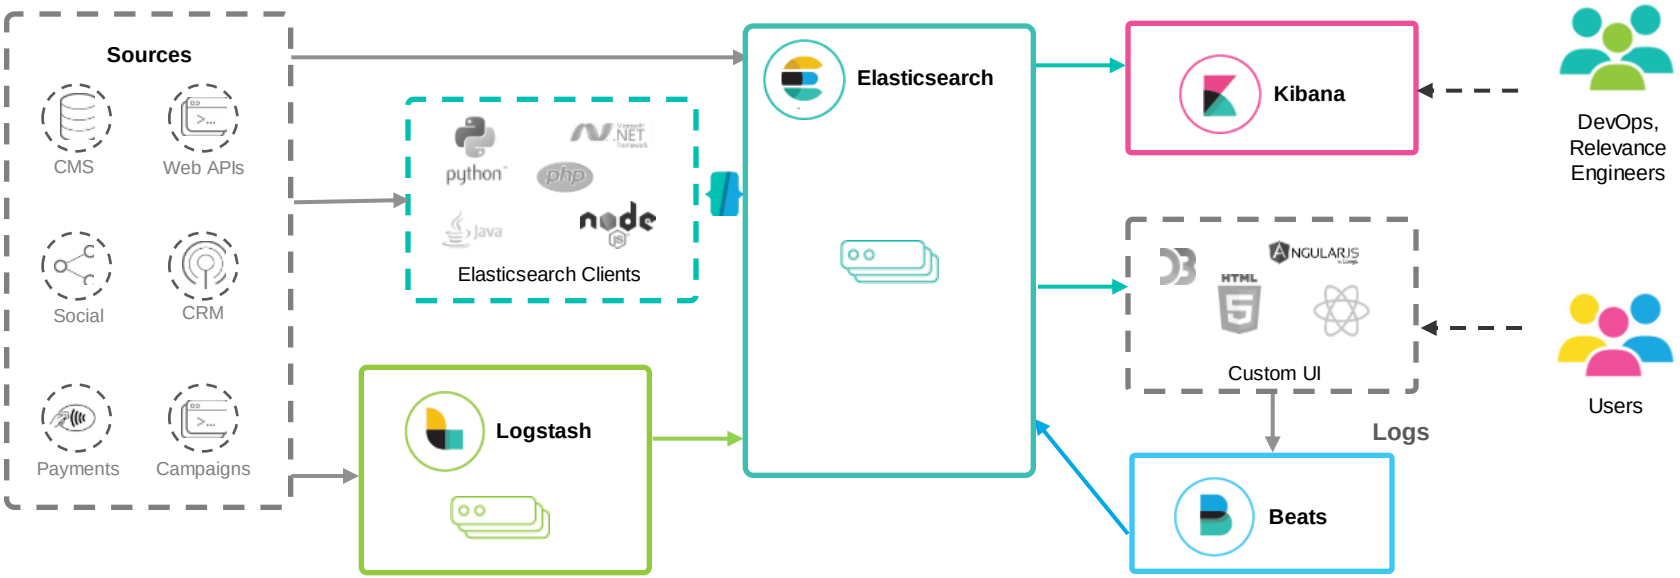

# Application or Knowledge Search: Scaled

*A dedicated cluster for logs/intelligence*

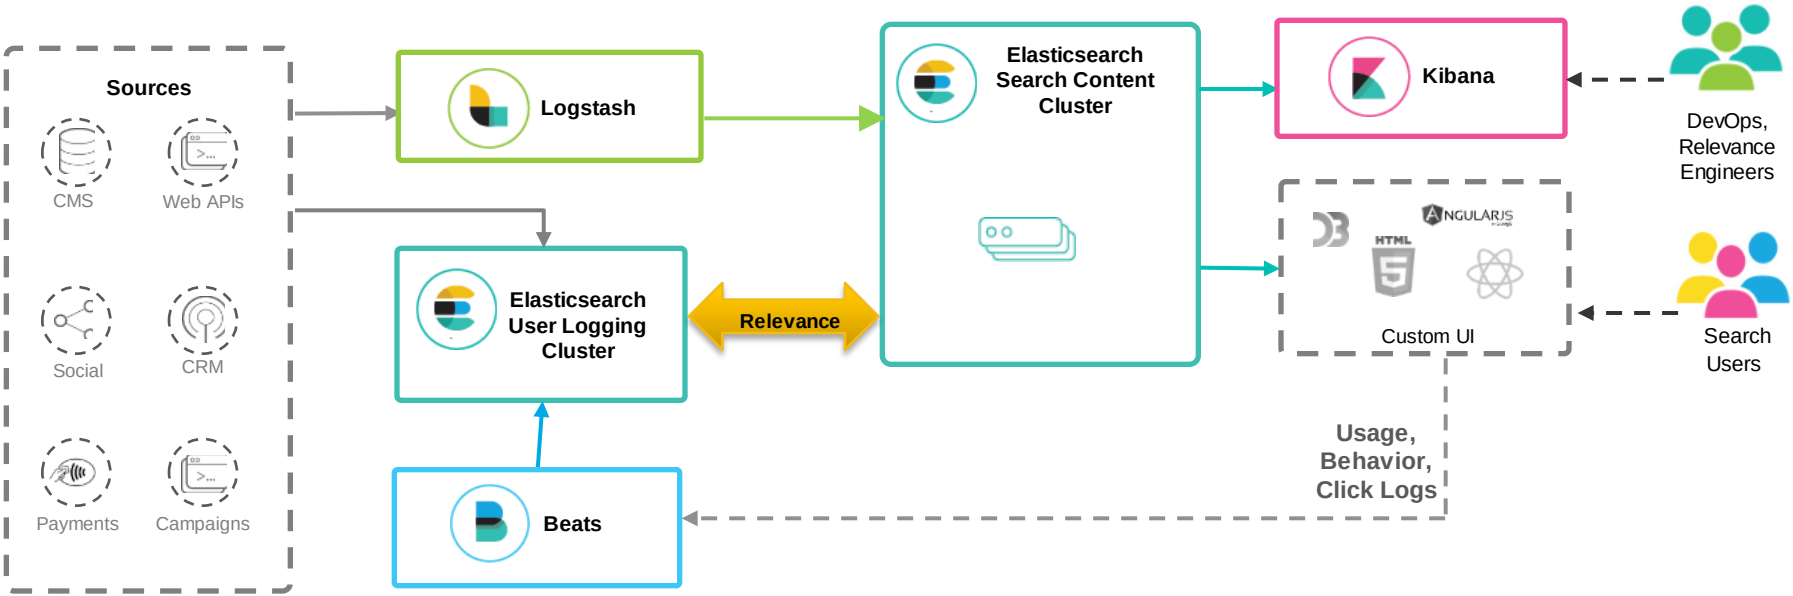

# Enterprise Search: Scaled

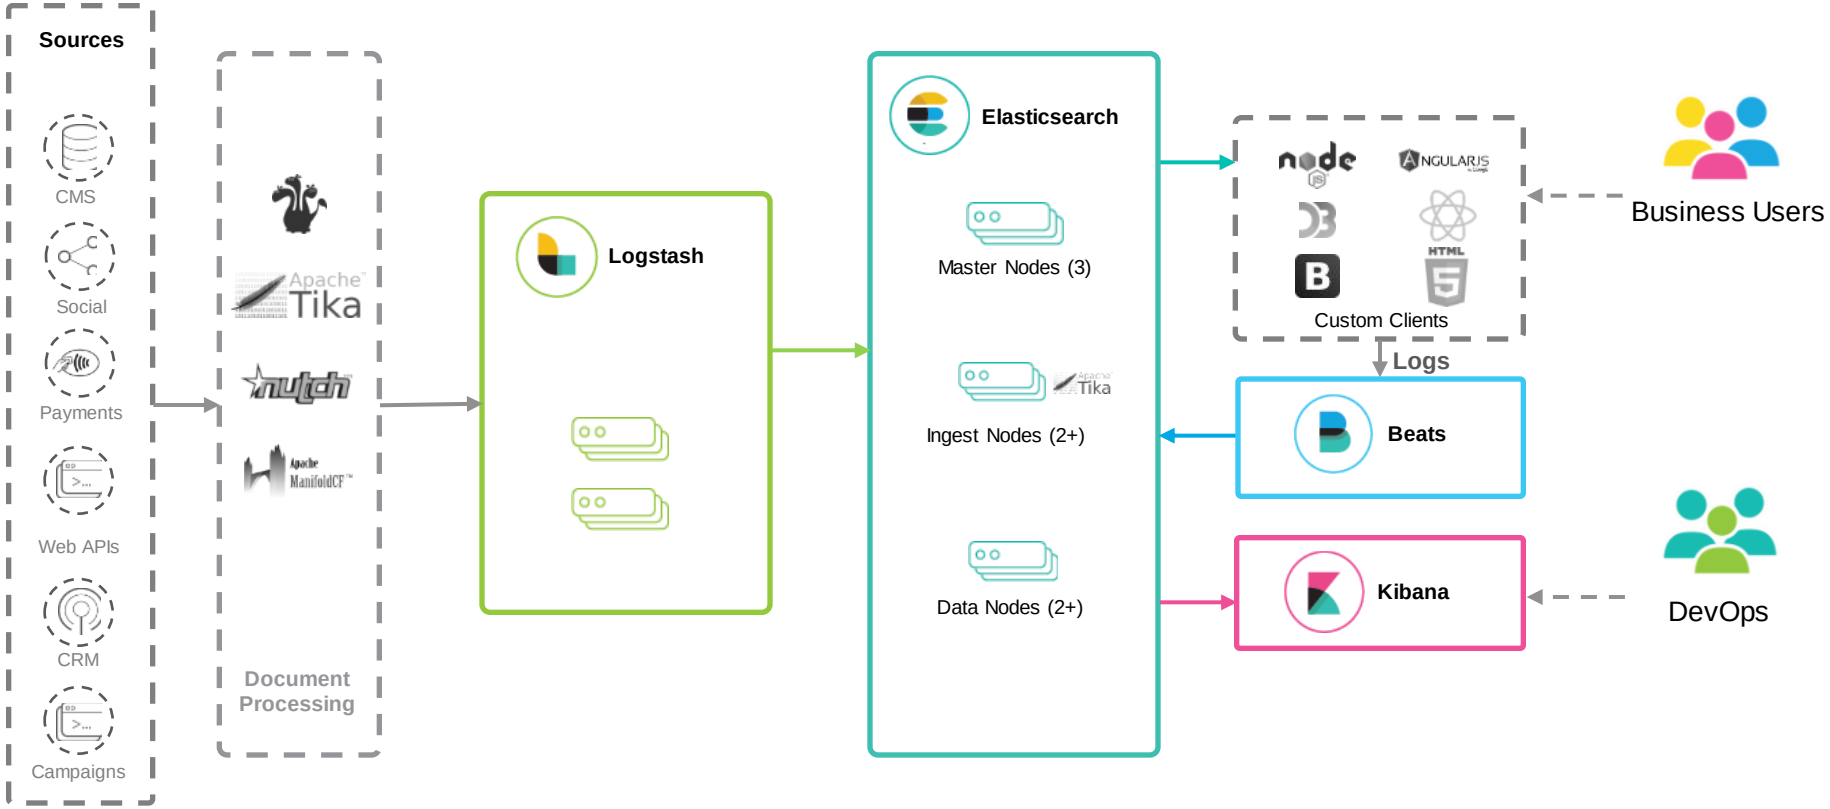

# Security Analytics Enterprise

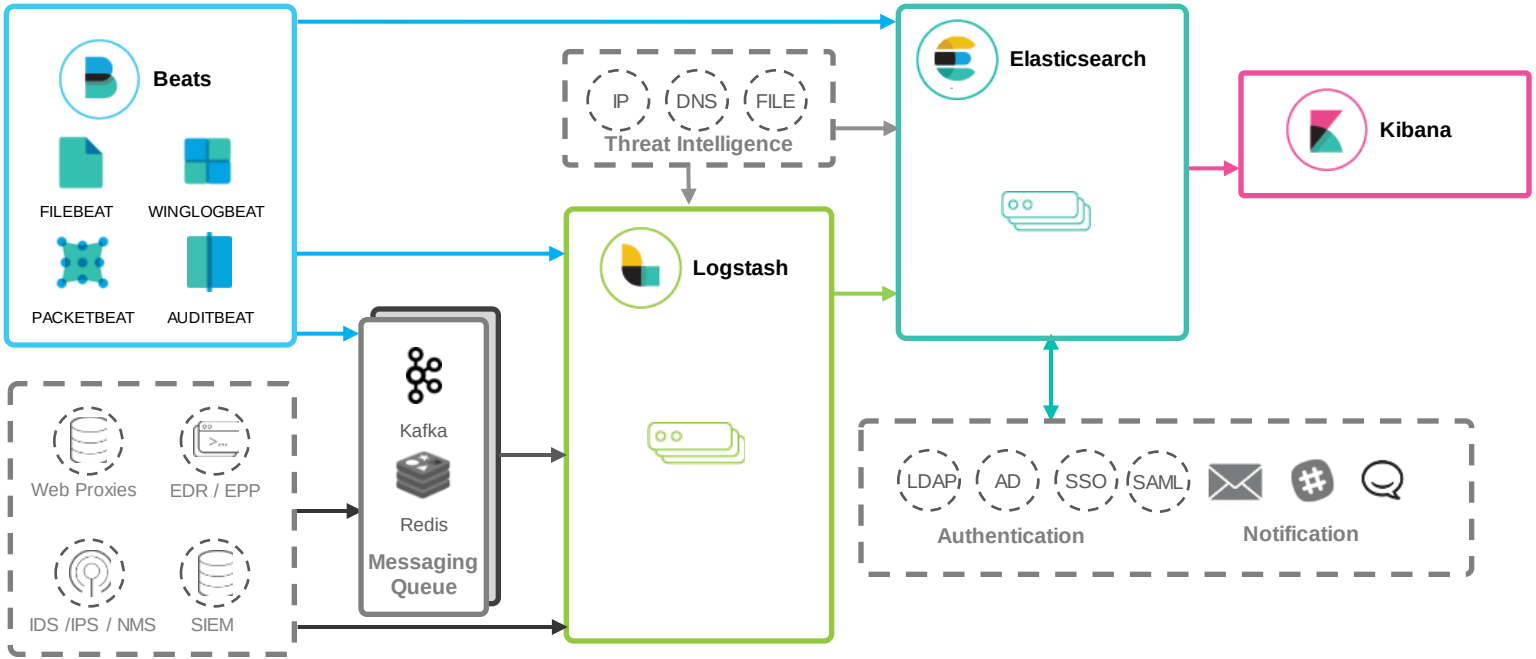

# Multi Data Center

# Multiple Data Centers, Duplicate Data

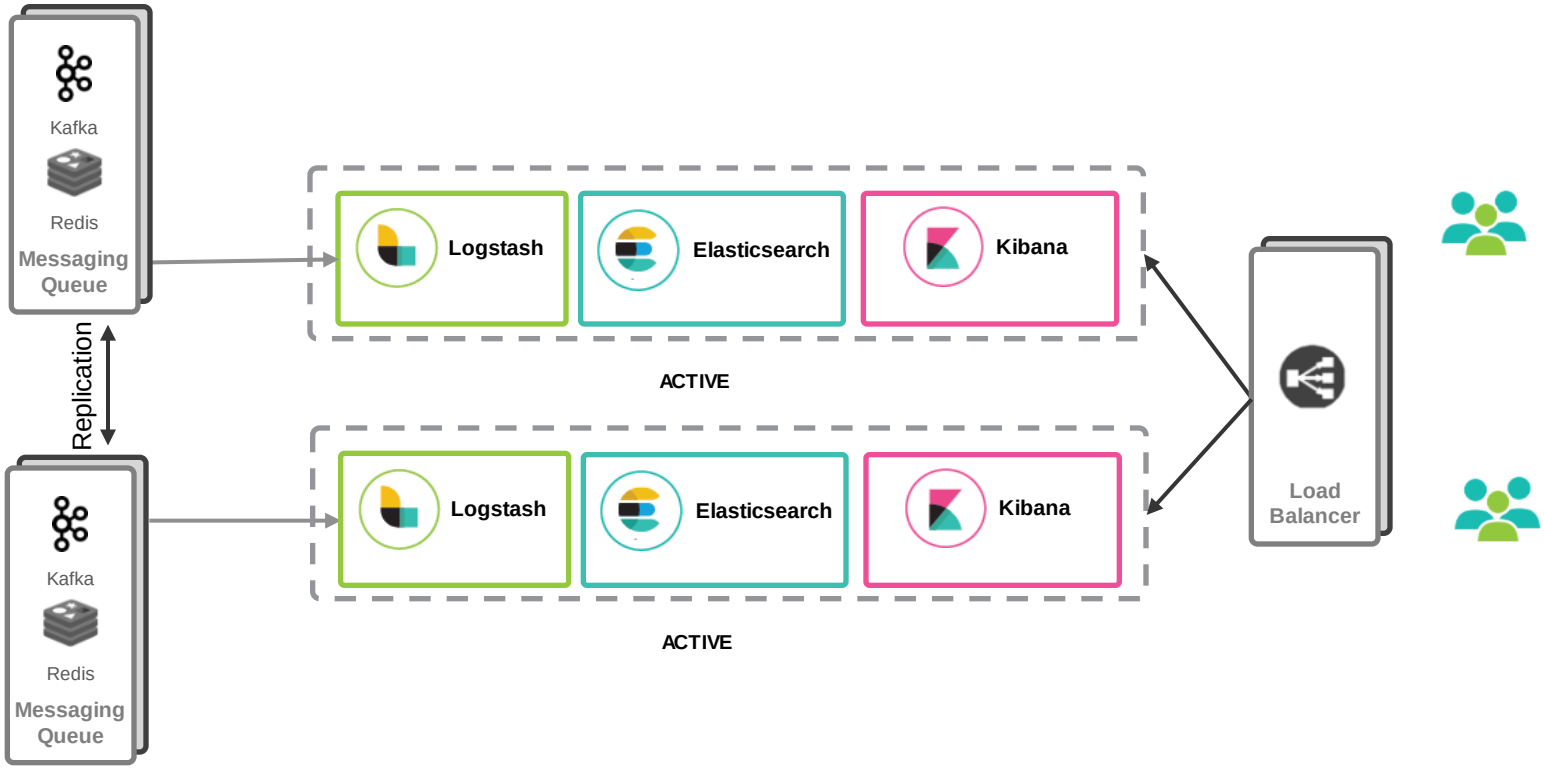

# Multi Data Centers with a Queue at Each DC

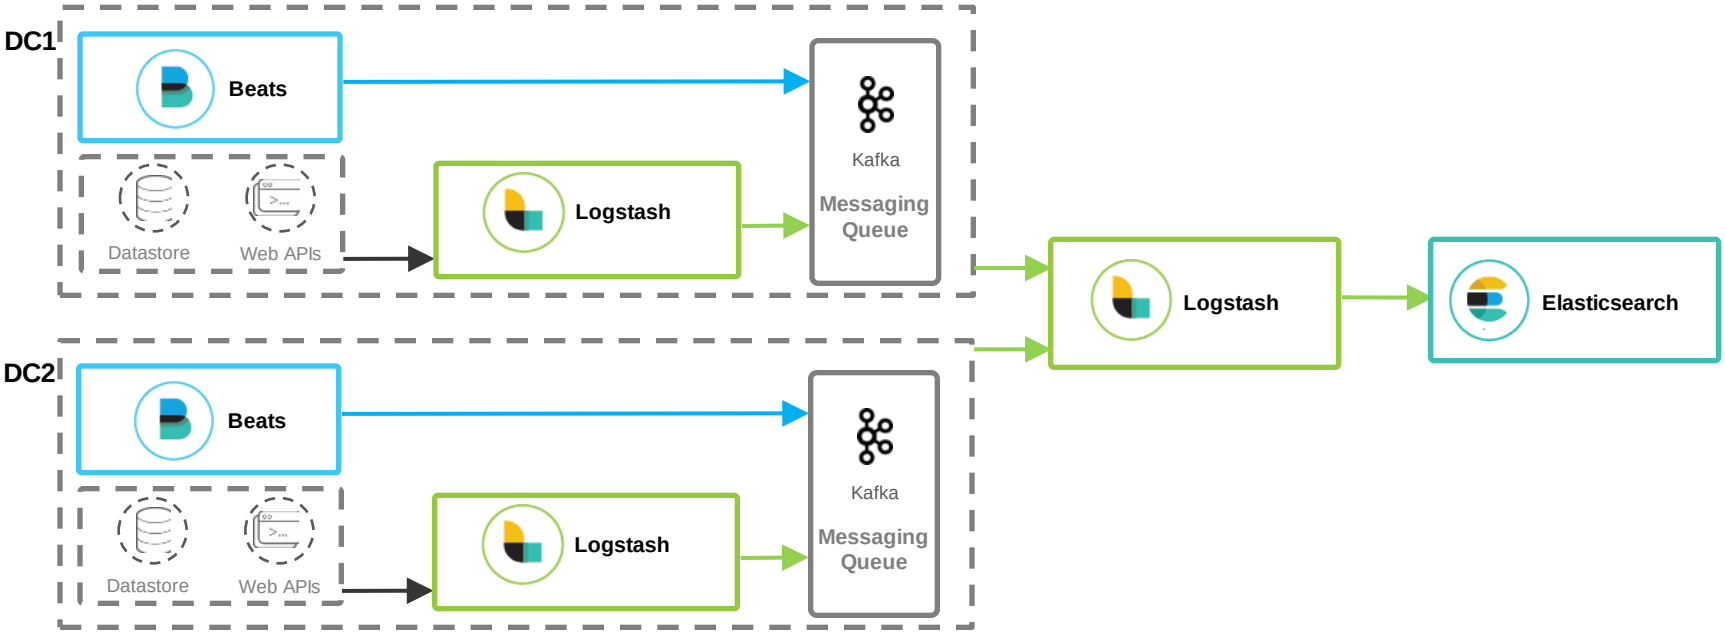




elastic

# Multi Data Center, Distinct Data and Cross-Cluster Search

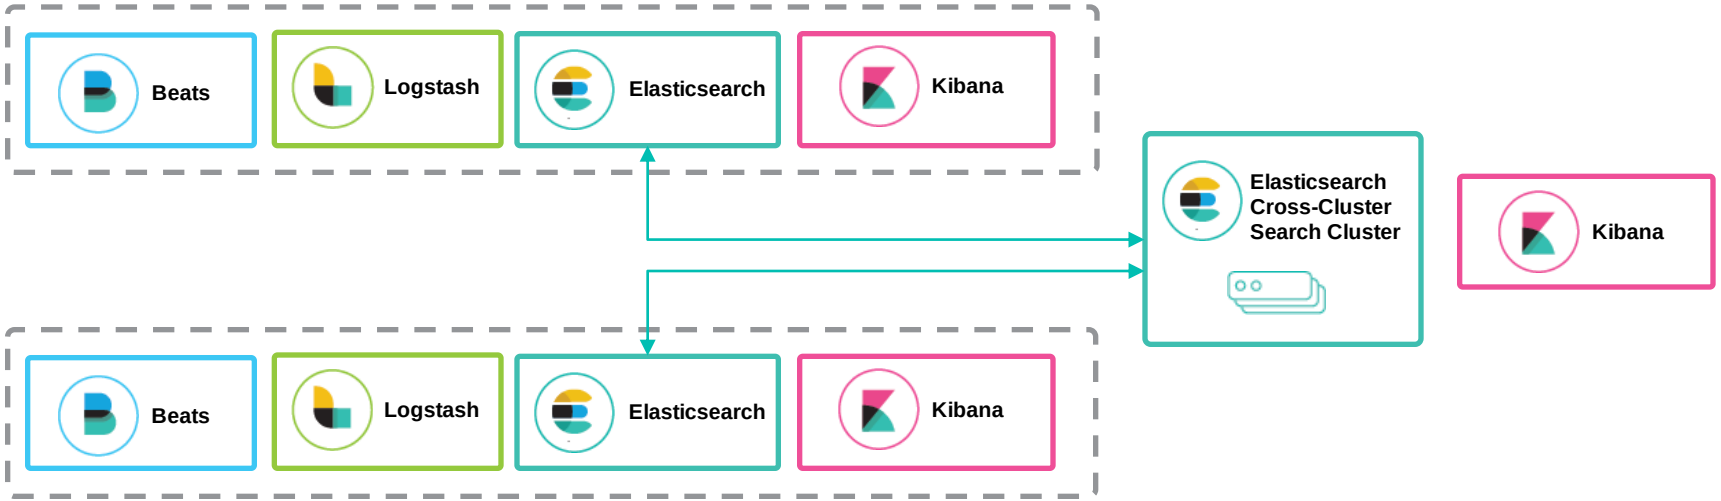

# Scaling Kibana

# High Availability

*Pair two coordinating nodes with two independent Kibana nodes*

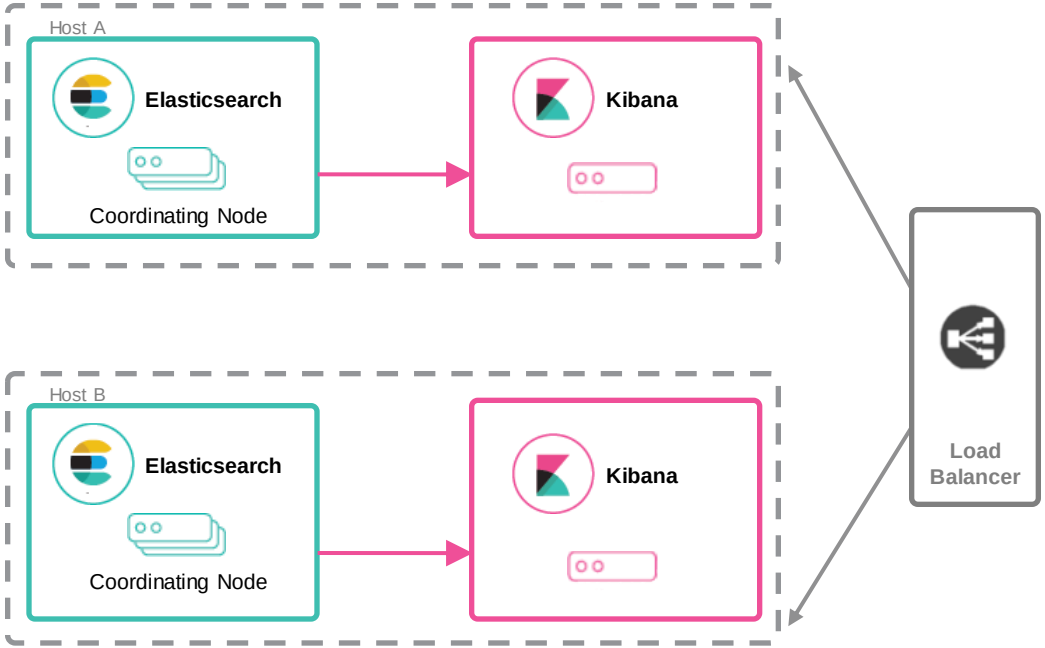




elastic

# Separating Content by Groups

*Isolate user content by group in different Kibana instances*

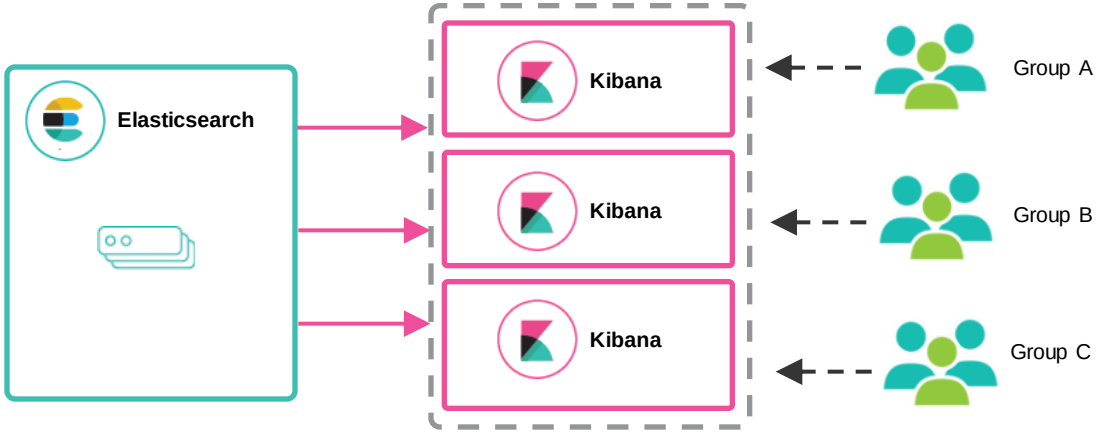




elastic

# Elastic Enterprise

## A SaaS that fits the enterprise

# Elastic Cloud Enterprise

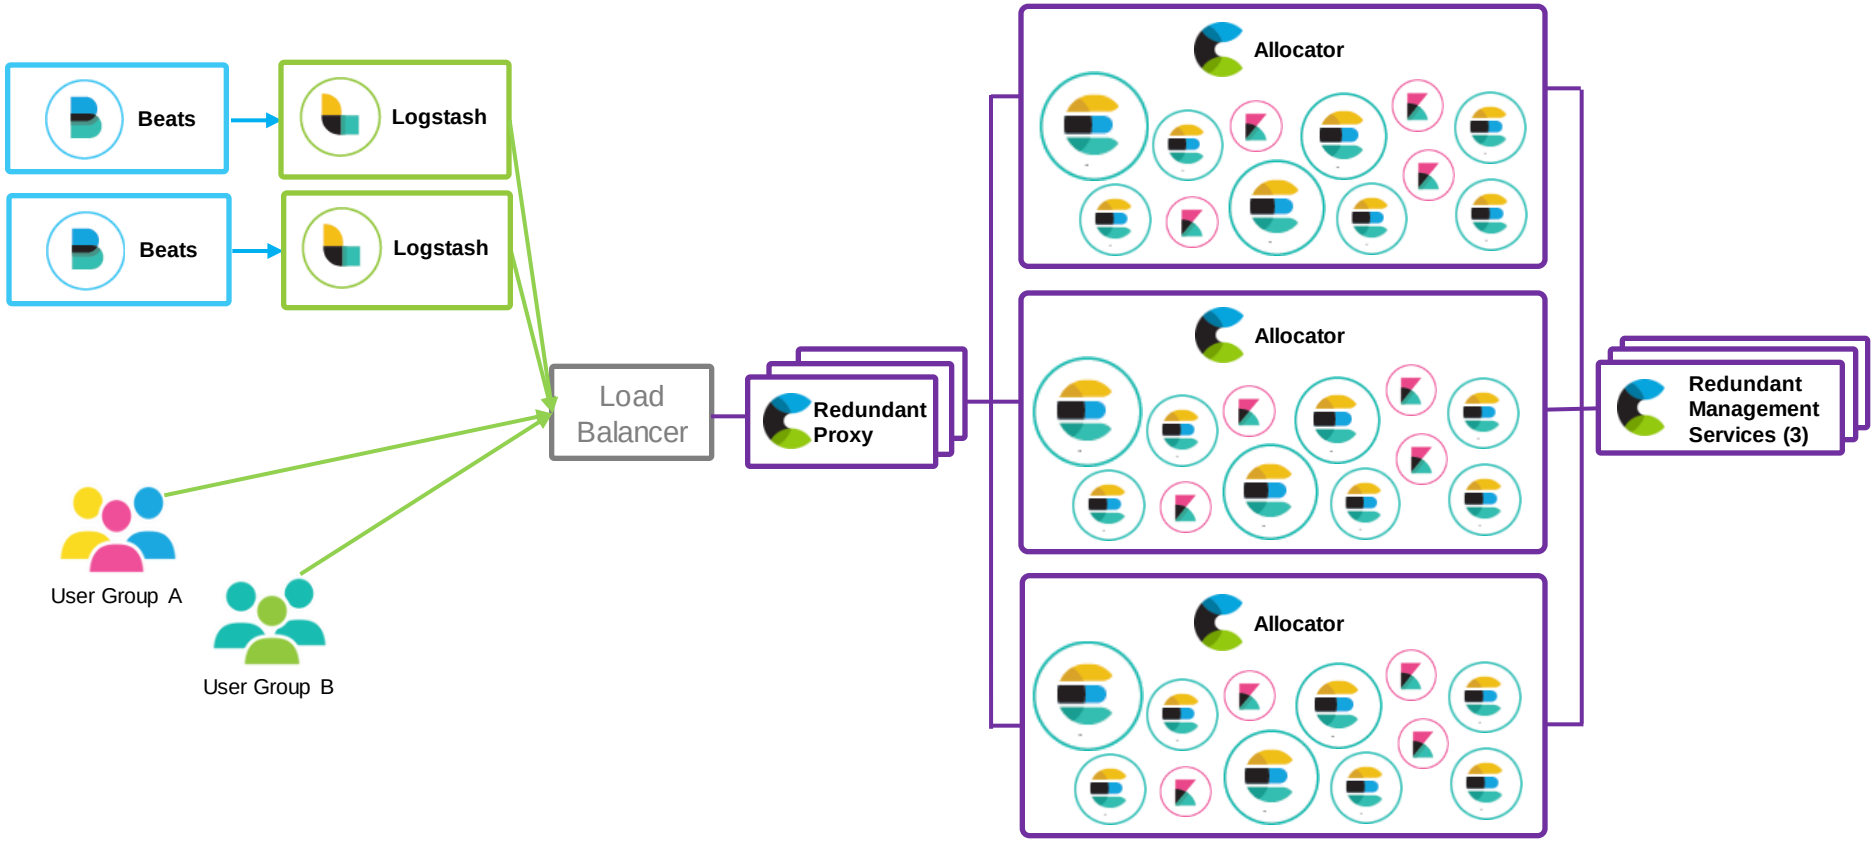

# Moving to Elastic Cloud Enterprise

*Options for easily moving your Elasticsearch cluster*

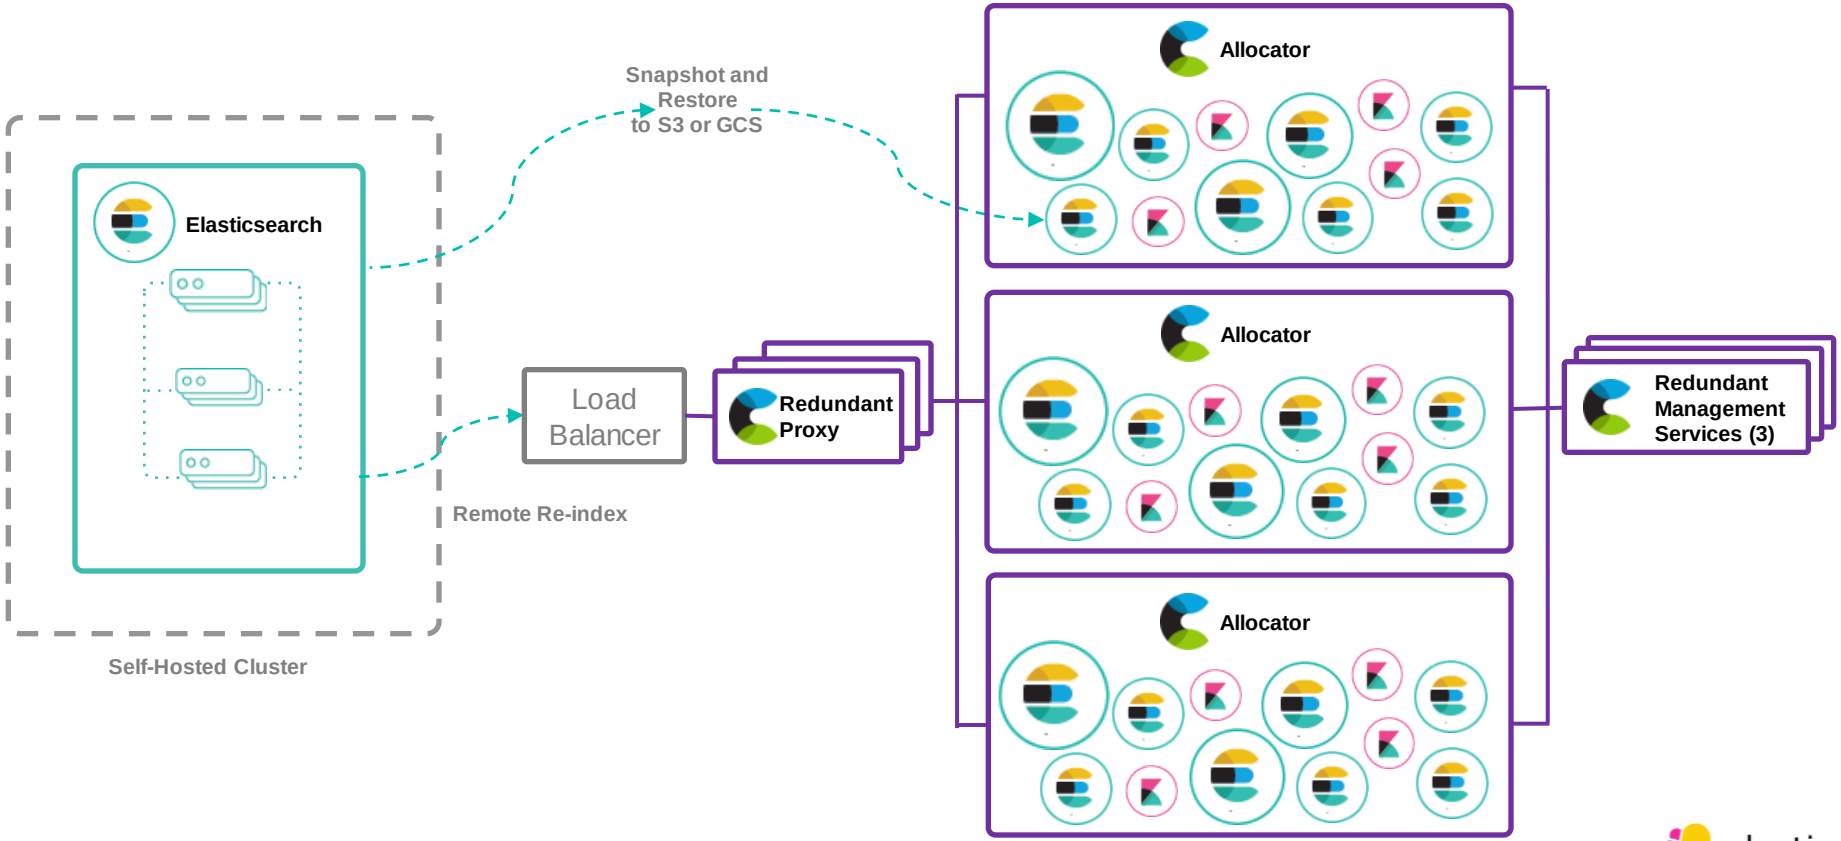






elastic

# Elastic Cloud Enterprise high level architecture

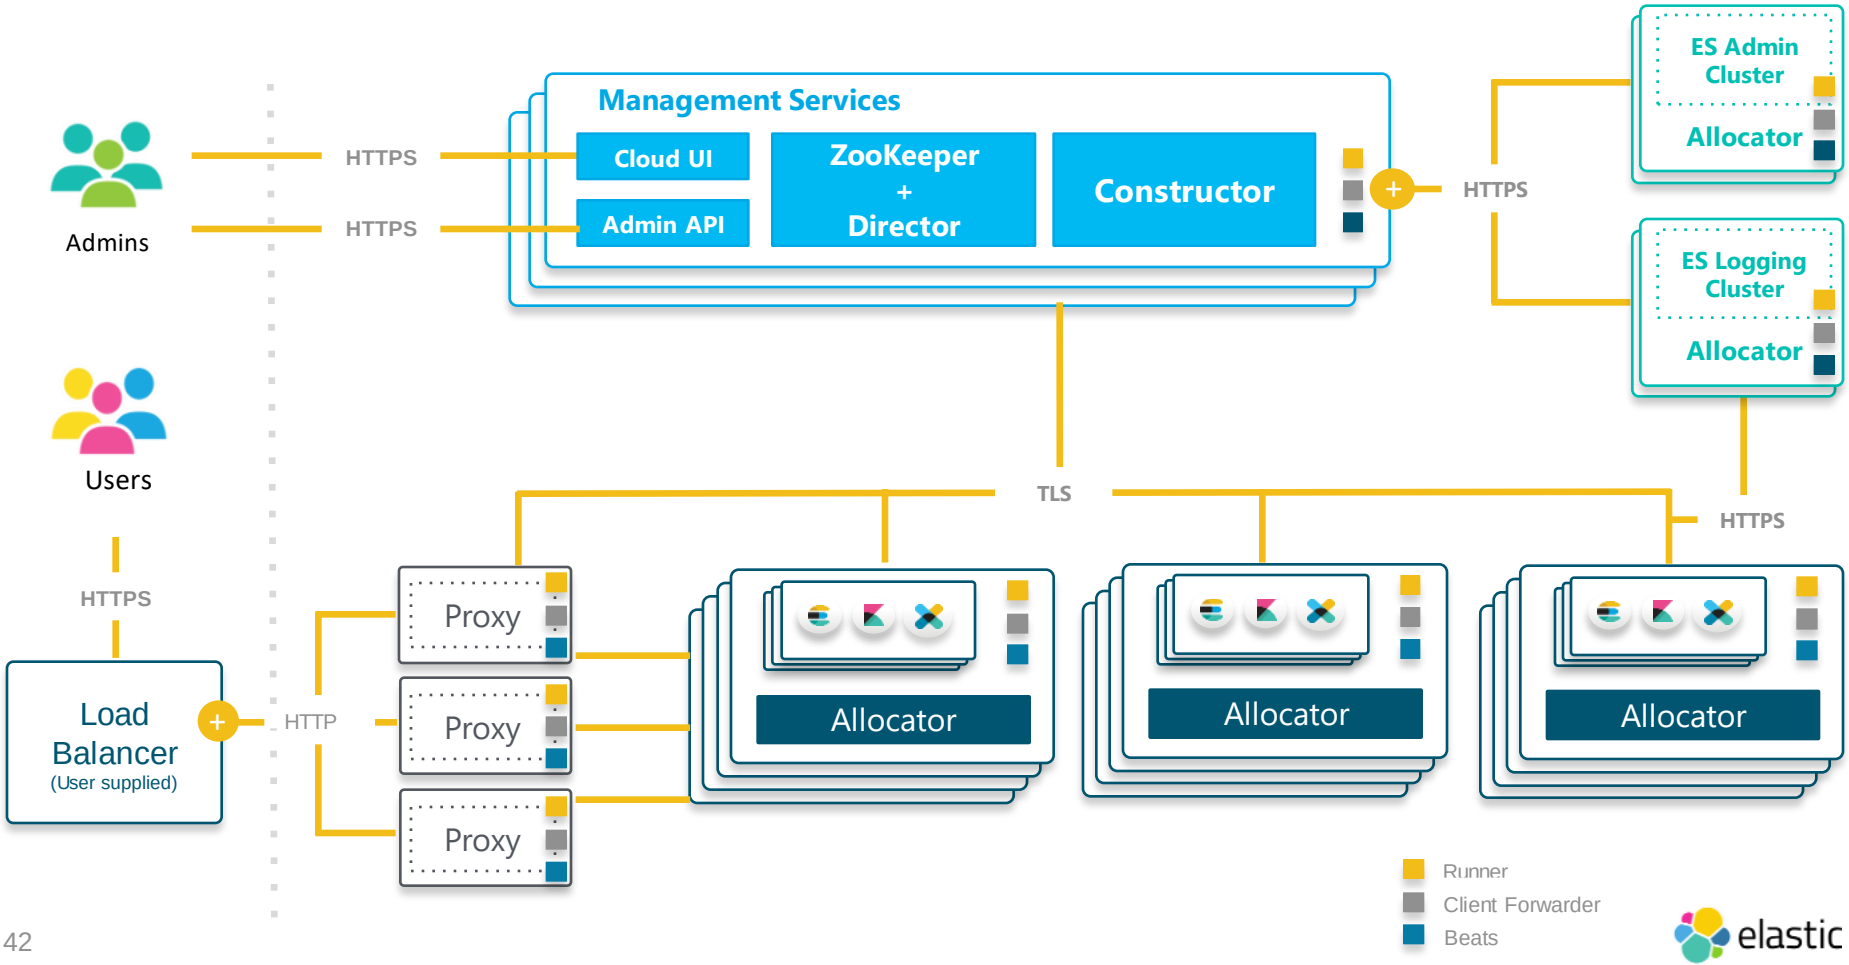

# Hardware

# Hardware Recommendations

*For On-Premise or IaaS Cloud Deployments*






** Data volumes are guidelines*

elastic

# Hardware Recommendations

## *For Elastic Cloud Enterprise Deployments*

- *Data volumes are guidelines*
- *Multi data tiers available from ECE 1.2*

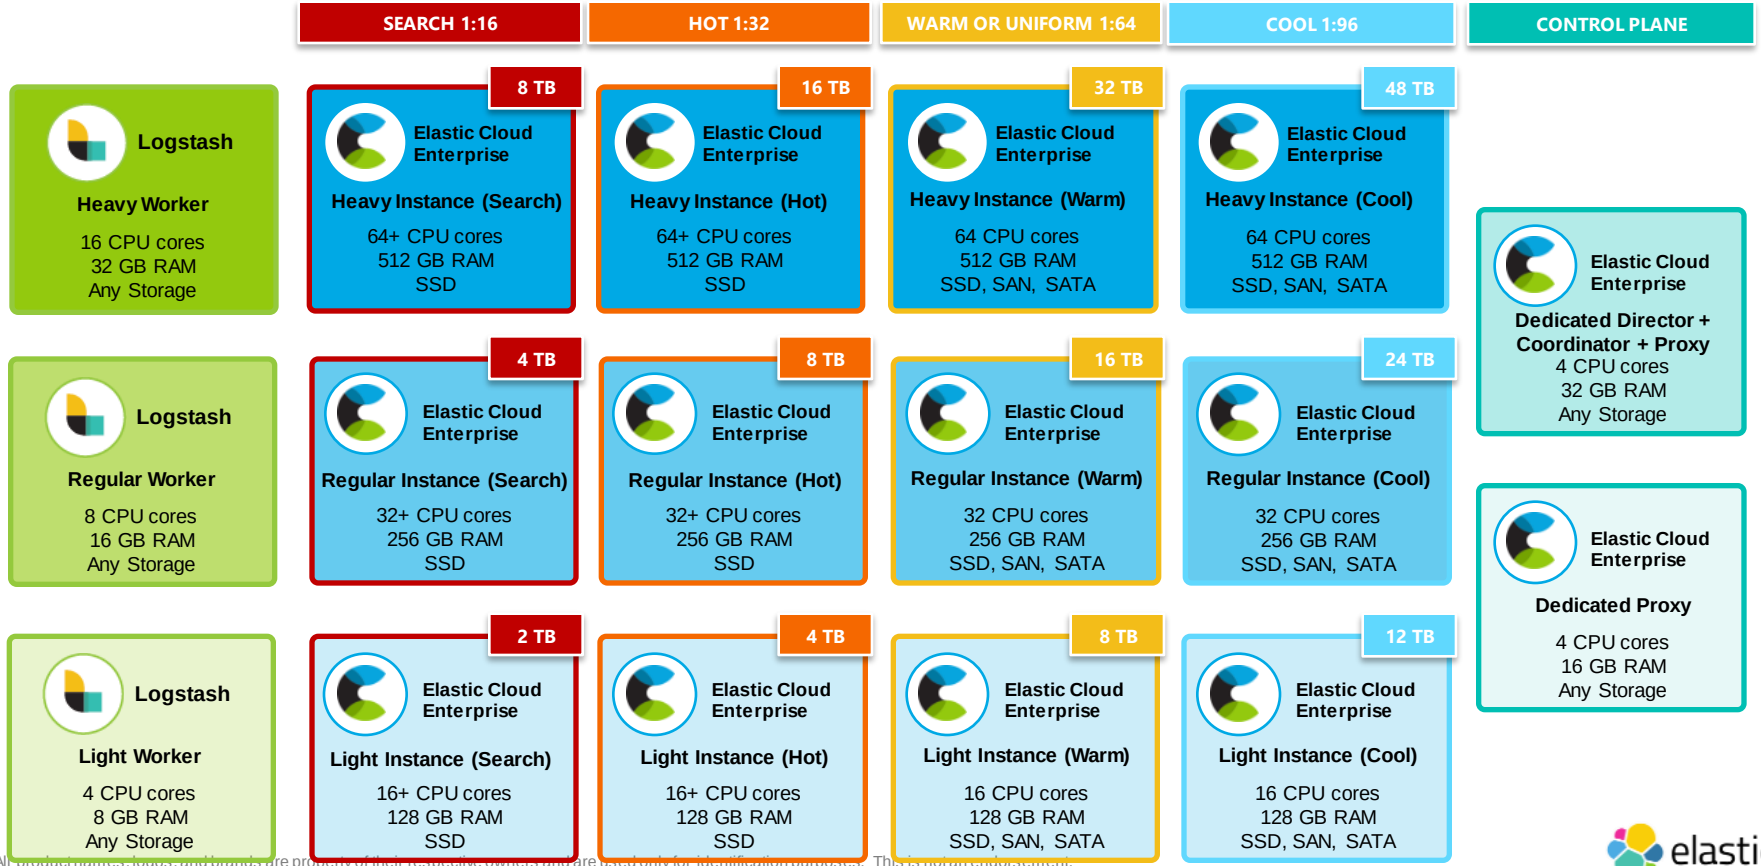

# Inside a Cluster

# Elasticsearch Node Types

*Nodes can play one or more roles, for workload isolation and scaling*

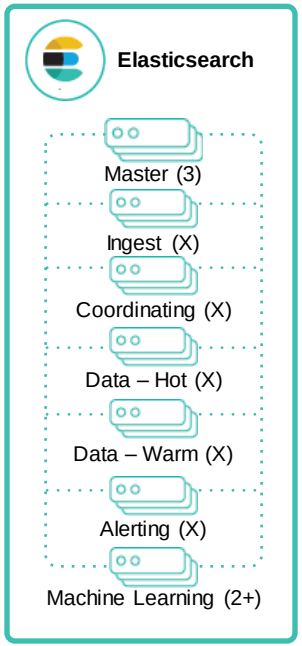


- Master Nodes
  - Control the cluster, requires a minimum of 3, one is active at any given time
- Data Nodes
  - Hold indexed data and perform data related operations
  - Differentiated Hot and Warm Data nodes can be used
- Ingest Nodes
  - Use ingest pipelines to transform and enrich before indexing
- Coordinating Nodes
  - Route requests, handle search reduce phase, distribute bulk indexing
  - All nodes function as coordinating nodes
- Alerting Nodes
  - Run alerting jobs
- Machine Learning Nodes
  - Run machine learning jobs

# Elasticsearch Coordinating Nodes Detail

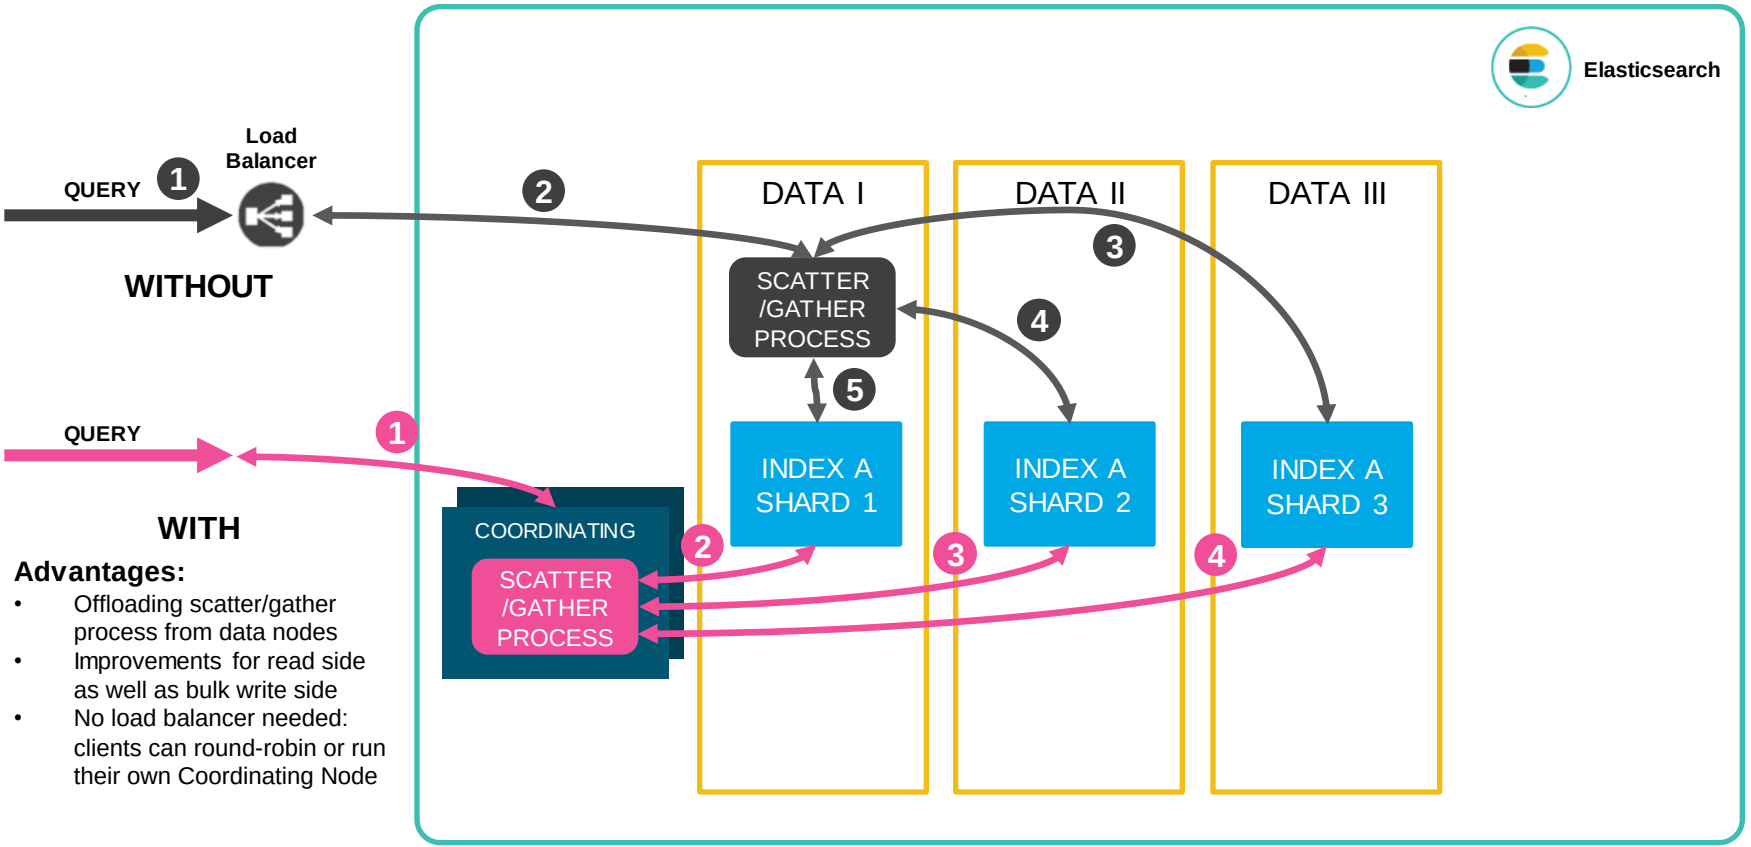


**Advantages:**

- Offloading scatter/gather process from data nodes
- Improvements for read side as well as bulk write side
- No load balancer needed: clients can round-robin or run their own Coordinating Node

# Elasticsearch Custom Routing

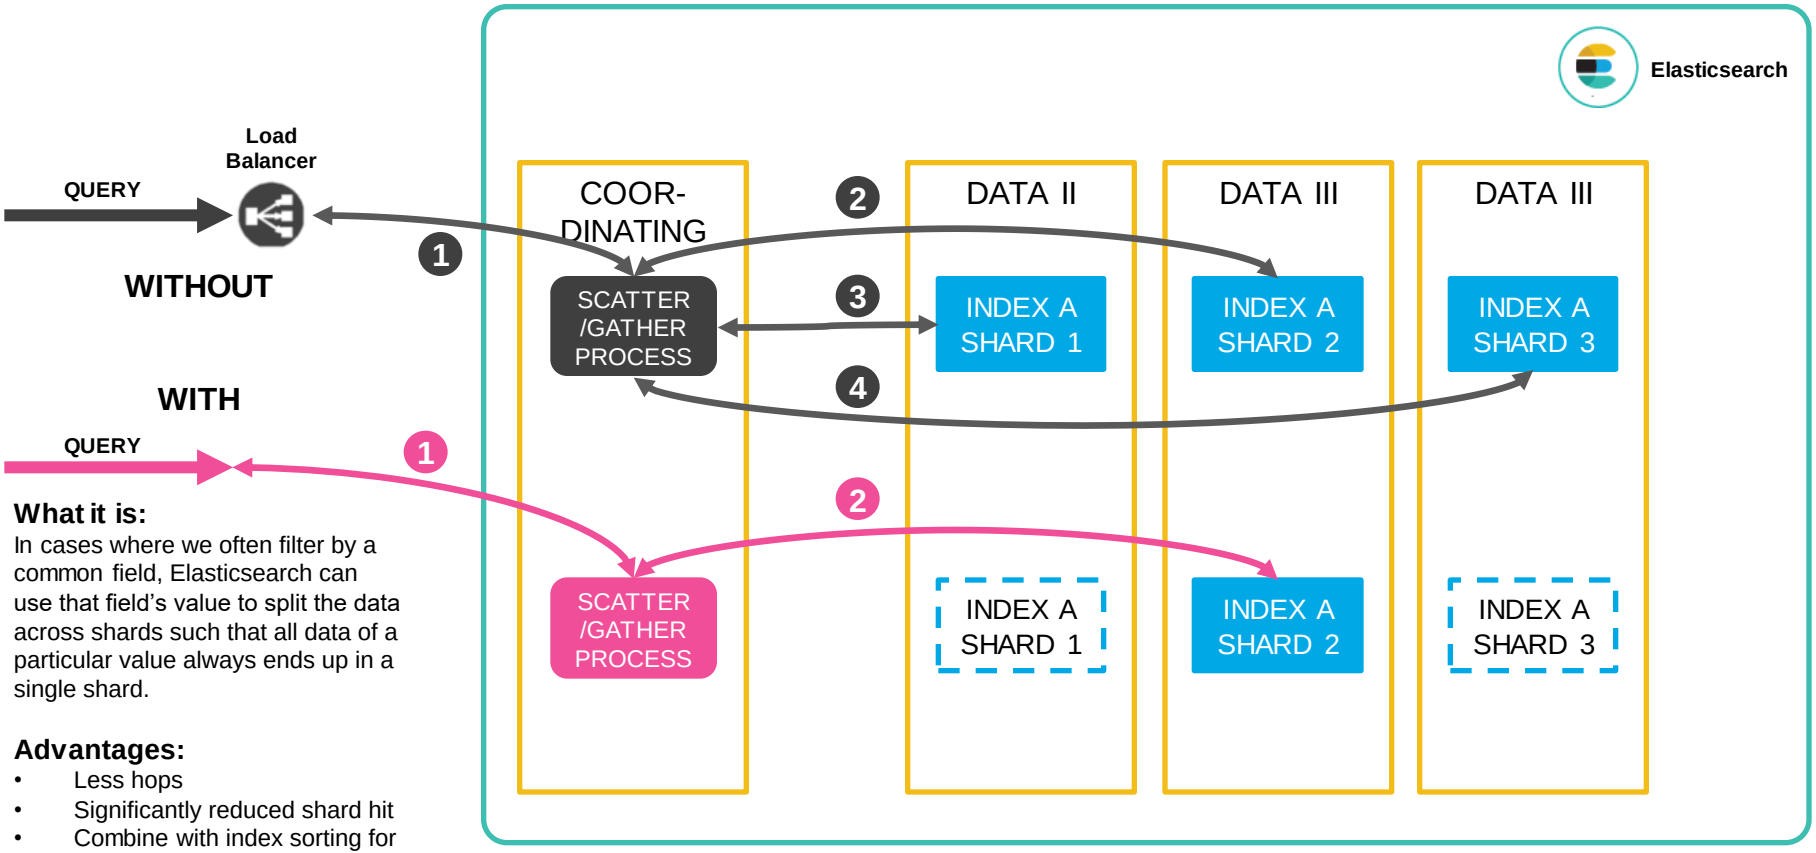


**What it is:**
In cases where we often filter by a common field, Elasticsearch can use that field's value to split the data across shards such that all data of a particular value always ends up in a single shard.

**Advantages:**

- Less hops
- Significantly reduced shard hit
- Combine with index sorting for crazy-fast searching